

# CS 288 Intensive Programming in Linux

Professor Ding, Xiaoning

**This content may NOT be uploaded, shared, or distributed, as it is protected.**

# Structures can organize data in different types

- Declared using struct with member types and names included in braces.
- struct variables can be declared with the struct.

```
struct transaction
{
    int id;
    float amount;
    char name[20];
    char addr[30];
} t, *pt;
```

```
struct transaction
{
    int id;
    float amount;
    char name[20];
    char addr[30];
};
struct transaction
t, *pt;
```

```
typedef struct
{
    int id;
    float amount;
    char name[20];
    char addr[30];
} transaction;
transaction t,*pt;
```

# Accessing struct members using . or ->

- A member in a **struct variable** can be access using .
- A member in a **struct pointed by a pointer** can be access using -> or by dereferencing the pointer first and then using .

```
pt=&t;  
printf("id: %d, name: %s, addr: %s",  
      t.id, pt->name, (*pt).addr);
```

# Pointer members in a struct

Some members need to have their memory dynamically allocated or location dynamically determined.


```
struct transaction{
    int id;
    float amount;
    char name[20];
    char *addr;
} t1, t2;
t1.addr=(char *)malloc(30);
t2.addr=another_str;
```

Extra pointers can be added to structs to support data structures, e.g., linked list, stack, queue, tree, graph, ...

```
struct transaction{
    int id;
    float amount;
    char name[20];
    char *addr;
    struct transaction *next;
} t1, t2;
t1.next = &t2;
```

# Linked Lists

- A linked list is a sequence of connected nodes.
- Each node contains at least
  - Some data
  - A pointer to the next node in the list
- The head pointer points to the first node
- The last node points to NULL
- The tail pointer (optional) points to the last node.




# Why linked list?

- Often the maximum size of the list cannot be estimated.
  - Static arrays have fixed sizes.
  - Extending dynamic arrays (malloc-ed mem space) may need to copy data from the old and smaller space to a larger new space.
- Usually there are updates in the middle of the list, e.g., insertion, deletion, re-arranging, etc. Overhead is high with arrays.
  - Inserting a new element in the front or deleting the first element requires shifting all the elements in the array
  - On average, half of the lists needs to be moved for insertion/deletion.
- Compared to an array, linked list uses only as much space as is needed (requires extra-space for pointers)

# Building a linked list

- **Declare node type** --- self-referential struct
- **Create nodes** --- allocate memory on-demand, initialize members
- **Link nodes to the list** --- find a location on the list (the previous and/or the next node) and update pointers in these nodes and the new node.
- **Keep the head pointer updated** --- If the address is lost, the whole list may be lost.
- **Ensure the next pointer of the last node to NULL.**
- **If there is a tail pointer, keep it updated**



```
#include <stdio.h>
#include <stdlib.h>
struct node{
    int id;
    char name[20];
    struct node *next;
};
```

Declare node type  
self\_referential struct

head/tail pointers always  
point to the first/last node.

Create and initialize  
a new node

```
int main() {
    struct node *head=NULL, *tail=NULL, *pnode;
    while(1) {
        pnode=(struct node *)malloc(sizeof(struct node));
        printf("id:"); scanf("%d", &(pnode->id));
        if(pnode->id<0) break;
        printf("name:"); scanf("%s", pnode->name);
        pnode->next=NULL; /*ensure next pointer of last node is NULL */
        if(head==NULL) head=pnode;
        if(tail!=NULL) tail->next=pnode;
        tail=pnode;
    }
```

Link the new node to the end




```
pnnode=head;
while (pnnode!=NULL) {
    printf("id: %d\t name:%s\n", pnnode->id, pnnode->name);
    pnnode=pnnode->next;
}
}
```

## Traverse a linked list

- Start from the head pointer.
- Proceed following the next pointers of nodes.
- Stop when the last node is reached.
  - Last node: next pointer is NULL, or pointed by the tail pointer.

# Adding a node to a list

adding to the front




```
b->next = head;  
head = b;
```

Order of the operations is important

```
head = b;  
b->next = ?
```

# Adding a node to a list

## adding to the front




```
b->next = head;  
head = b;
```

Order of the operations is important

```
head = b;  
b->next = ?
```


## Inserting into the middle




```
b->next = a->next;  
a->next = b;
```

Consider `a->next` as a "head pointer" to the rest of the list.

# Deleting a node



```
p = a->next;  
a->next = a->next->next;  
...  
free(p);
```



```

p = a->next;
a->next = a->next->next;
...
free(p);

```

## Deleting the first node

```

p = head;
head = head->next;
...
free(p);

```

## Deleting the last node

?

# Inserting/deleting a node: what to notice?

- Must handle different scenarios in different ways
  - E.g., for adding a node, check whether the new node is to be added to an empty list, to the front, the middle, or the end.
- Pay special attention to whether your code may dereference a NULL pointer.
  - Segmentation faults.
- Avoid losing the pointer pointing to a node
  - It becomes "inaccessible" if there are no pointers points to it.
- Keep the head pointer and the tail pointer (if you have one) updated.

# Insertion sort

```
#include <stdio.h>
#include <stdlib.h>

struct node{
    int data;
    struct node *next;
};

void show_list(struct node *head) {
    struct node *p=head;
    if(head==NULL) return;
    while(p!=NULL){
        printf("->%d", p->data);
        p=p->next;
    }
    printf("\n");
}
```

```
struct node *add_to_sorted_list(
    struct node *head,
    struct node *p){
    struct node *prev_node, *curr_node;

    if(head==NULL) return p;
    if(p->data < head->data){
        p->next=head;
        return p;
    }
    prev_node=head;
    curr_node=prev_node->next;
    while(curr_node != NULL){
        if(p->data < curr_node->data) break;
        prev_node=curr_node;
        curr_node=curr_node->next;
    }
```

# Insertion sort

```
    prev_node->next=p;
    p->next=curr_node;
    return head;
}
int main(){
    int i;
    struct node *p, *head=NULL;
    while(scanf("%d",&i)!=EOF){
        p=(struct node *)malloc(sizeof(struct node));
        p->data=i;
        p->next=NULL;
        head=add_to_sorted_list(head,p);
    }
    show_list(head);
}
```




# Using a dummy node to unify different scenarios

- With a dummy in the front, every node appear to be added after the next field of a node.
- Some linked list implementations keep the dummy node as a permanent part of the list.
  - Every list has a dummy node at its head.
  - The empty list is not represented by a NULL pointer.



```
struct node *dummy=(struct node *)
                malloc(sizeof(struct node));
struct node *add_to_sorted_list(
                struct node *head,
                struct node *p){
    struct node *prev_node, *curr_node;
    dummy->next=head;
    head=dummy;
    prev_node=head;
    curr_node=prev_node->next;
    while(curr_node != NULL) {
        if(p->data < curr_node->data) break;
        prev_node=curr_node;
        curr_node=curr_node->next;
    }
    prev_node->next=p;
    p->next=curr_node;
    return dummy->next;
```

# Using pointer to pointer to reduce # of scenarios

node \* head



node \* p->next;



node \* p

```
struct node* add_to_sorted_list(struct node *head,  
                                struct node *p){
```

```
    struct node **curr, *next;
```

```
    if(head==NULL) return p;
```

```
    curr=&head;
```

```
    do{
```

```
        if(p->data < (*curr)->data) break;
```

```
        curr=&((*curr)->next);
```

```
    }while(*curr!=NULL);
```

```
    next=*curr;
```

```
    *curr=p;
```

```
    p->next=next;
```

```
    return head;
```

```
}
```

# Debugging linked list programs

## Segmentation fault (segfault)

- caused by a program accessing the memory it should not
- Use after free, Dereferencing NULL pointers
- Use gdb
  - Run program in gdb
  - Execution stop on segfault
  - Check stack using backtrace (bt) or where command.
  - check the values of related variables and pointers.

```
struct node *add_to_sorted_list(
    struct node *head, struct node *p){
    struct node *prev_node, *curr_node;
    if(head==NULL) return p;
    if(p->data < head->data){
        p->next=head;
        return p;
    }
    prev_node=head;
    curr_node=prev_node->next;
    while (p->data < curr_node->data){
        prev_node=curr_node;
        curr_node=curr_node->next;
    }
    prev_node->next=p;
    p->next=curr_node;
    return head;
}
```

# Debugging linked list programs

Show linked list every time after the list is changed.

```
struct node *add_to_sorted_list(
    struct node *head, struct node *p){
    struct node *prev_node, *curr_node;

    if(head==NULL) return p;
    if(p->data < head->data){
        p->next=head;
        return p;
    }

    prev=head;
    curr_node=prev_node->next;
    while(curr_node != NULL) {
        if(curr_node->data > < p->data)
            break;
        prev_node=curr_node;
        curr_node=curr_node->next;
    }
```

```
        prev_node->next=p;
        p->next=curr_node;
        return head;
    }
    int main(){
        int i;
        struct node *p, *head=NULL;
        while(scanf("%d",&i)!=EOF){
            p=(struct node *)malloc(
                sizeof(struct node));


            p->data=i;
            p->next=NULL;
            head=add_to_sorted_list(head,p);
            show_linkedlist(head);
        }
        show_linkedlist(head);
    }
```

1  
->1  
3  
->1->3  
5  
->1->5->3  
7  
->1->7->5->3


# Other types of linked list

## Double linked list

```
struct node{  
    int id;  
    char name[20];  
    struct node *next;  
    struct node *prev;  
};
```



To be covered in data structure courses, including other data structures, such as trees and graphs.



# CS 288 Intensive Programming in Linux

Professor Ding, Xiaoning

# Programs, processes, and threads

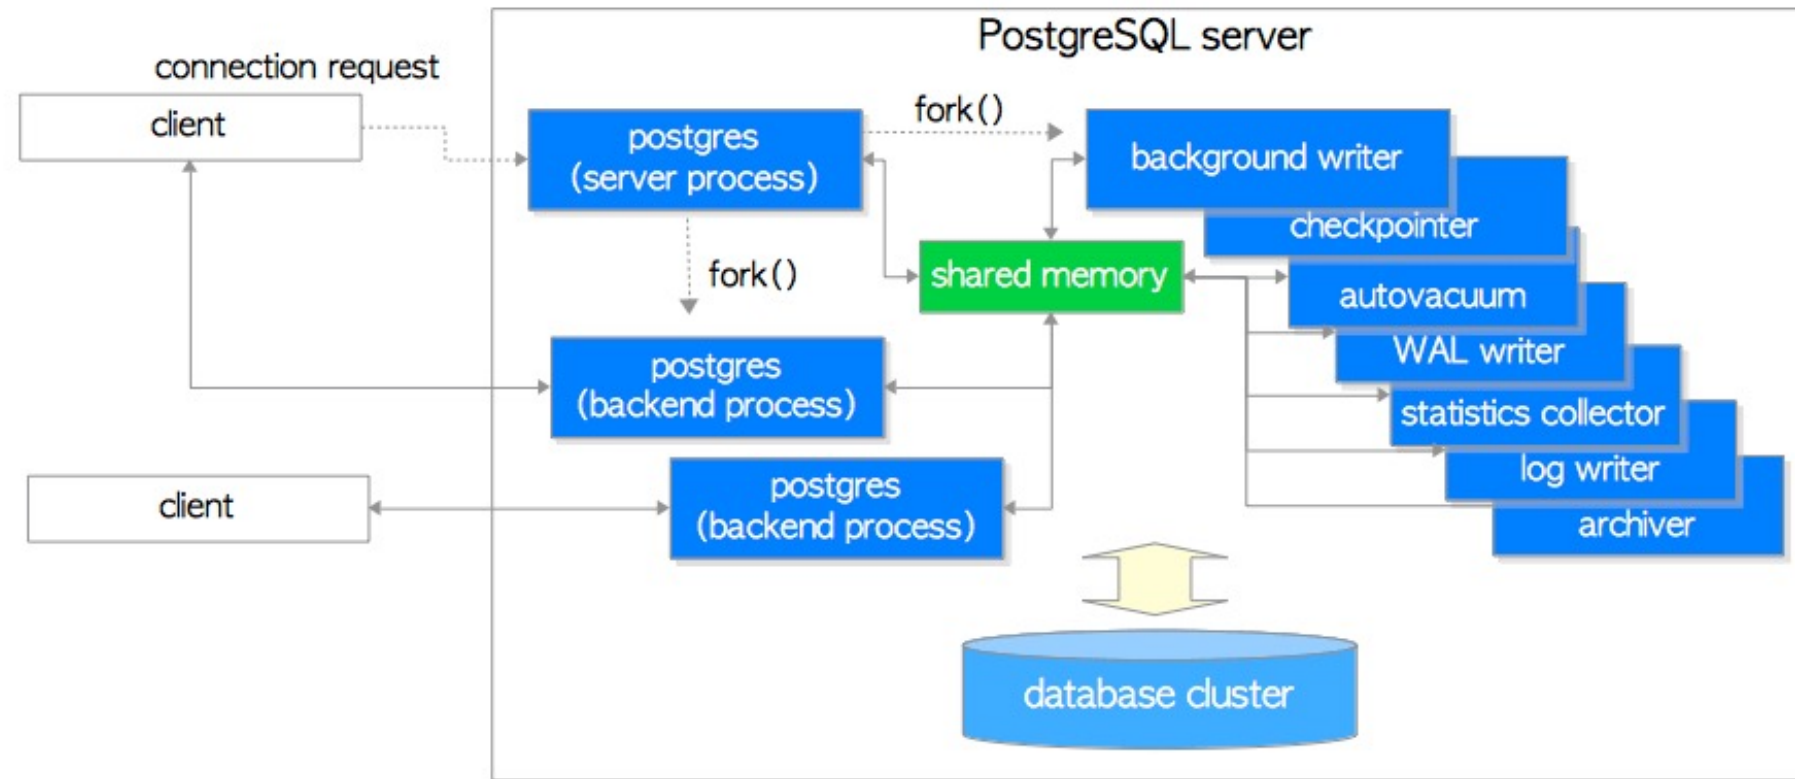
- Program: source code or executable file
- Process: running program
  - A binary image loaded into memory
  - An instance of virtual memory
    - Heap and stacks
  - kernel resources and data structures
    - Records about which files have been opened, memory that has been allocated, etc.
    - a unique process ID number(PID)
  - One or more threads
- Thread: a unit of activity inside of a process, corresponding to a virtualized processor
  - a stack
  - processor state such as registers and an instruction pointer
  - Threads in a process share the same virtual memory space

# The use of multiple threads, multiple processes, and multiple programs

- The use of multiple threads (multi-threading) and/or multiple processes
  - Objectives
    - To enhance parallel processing
    - To increase response to the user
    - To utilize the idle time of the CPU (e.g., scheduling another thread/process when one thread/process is waiting for I/O)
  - Multiple threads or multiple processes are created in the same program.
  - Threads vs. processes: a trade-off between isolation and data sharing
- The use of multiple programs
  - Each program is a process.
  - To fit the need of running multiple tools and different tools developed in different programs/developers.
  - Most common case: shell



# Postgresql Internals

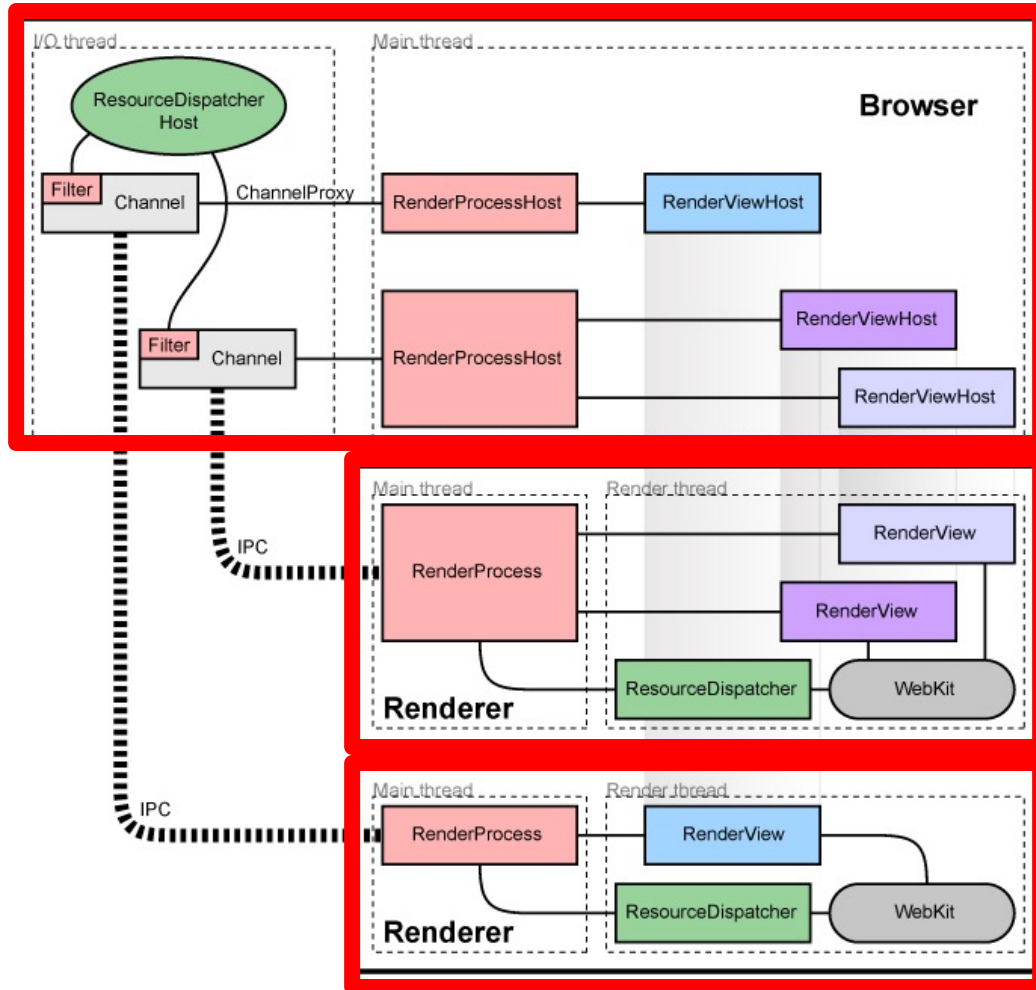


- A **postgres server process** is a parent of all processes related to a database cluster management.
- Each **backend process** handles all queries and statements issued by a connected client.
- Various **background processes** perform processes of each feature (e.g., VACUUM and CHECKPOINT processes) for database management.

# MySQL implements concurrent connections by spawning a thread-per-connection.

- relatively low overhead;
  - spawning a new thread occupies less memory than forking a new process in term of both time and memory.
  - Synchronization and data exchange incur much less cost between threads than between processes.
  - Context switches is less costly between threads than between processes
- It's not uncommon to scale MySQL to 10,000 or so concurrent connections.
- significant problems when scaling PostgreSQL past a few hundred active connect.

# Google Chrome architecture



- Separate processes for browser tabs to protect the overall application from bugs and glitches in the rendering engine
- Restricted access from each rendering engine process to others and to the rest of the system

<https://sites.google.com/a/chromium.org/dev/developers/design-documents/multi-process-architecture>

POSIX thread: pthread

# Pthreads --- POSIX thread library.

- POSIX (Portable Operating System Interface)
  - an IEEE standard that sets how UNIX systems look, act, and feel.
- The POSIX threads API is available on almost all UNIX-like operating systems.
  - if you write parallel code using Pthreads on a Linux machine, it will likely work on other UNIX variants.
- For C programs on Linux, pthreads API is implemented in Native POSIX Thread Library (NPTL).
- Need to link pthread library when compiling the source code.

**gcc mypthreadprog.c -pthread -o mypthreadprog**

```

#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

/* The "thread function" passed to pthread_create.  Each thread
 * executes this function and terminates when it returns from
 * this function. */
void *HelloWorld(void *id) {
    /* We know the argument is a pointer to a long, so we cast it
     * from a generic (void *) to a (long *). */
    long *myid = (long *) id;

    printf("Hello world! I am thread %ld\n", *myid);

    return NULL; // We don't need our threads to return anything.
}

int main(int argc, char **argv) {
    int i;
    int nthreads; //number of threads
    pthread_t *thread_array; //pointer to future thread array
    long *thread_ids;

```

```

// Read the number of threads to create from the command line.
if (argc !=2) {
    fprintf(stderr, "usage: %s <n>\n", argv[0]);
    fprintf(stderr, "where <n> is the number of threads\n");
    return 1;
}
nthreads = strtol(argv[1], NULL, 10);

// Allocate space for thread structs and identifiers.
thread_array = malloc(nthreads * sizeof(pthread_t));
thread_ids = malloc(nthreads * sizeof(long));

// Assign each thread an ID and create all the threads.
for (i = 0; i < nthreads; i++) {
    thread_ids[i] = i;
    pthread_create(&thread_array[i], NULL, HelloWorld, &thread_ids[i]);
}

/* Join all the threads. Main will pause in this loop until all threads
 * have returned from the thread function. */
for (i = 0; i < nthreads; i++) pthread_join(thread_array[i], NULL);

free(thread_array);      free(thread_ids);      return 0;
}

```

# Creating and using worker threads

- Include header file **pthread.h**
- Implement a **thread function**.
  - analogous to a `main()` function for a worker (created) thread. — a thread begins execution at the start of its thread function and terminates when it reaches the end.
  - Each thread executes the thread function using its private execution state, including its own stack and registers.
  - Input and output: **void \*** --- Addresses where real input data and output data are saved.
- Main thread (`main()` function) creates and “joins” worker threads.
  - The program starts as a single-threaded process --- main thread
  - Create worker threads: **pthread\_create()** function
  - Join worker threads: **pthread\_join()** function



# Creating a thread - pthread\_create()

```
int pthread_create (pthread_t *thread, pthread_attr_t *attr,  
void * thread_function, void * thread_args);
```

- pthread\_t \***thread**: address of a thread struct --- a handle to the thread.
- **attr**: thread attributes. Usually set to NULL to use default thread attributes.
- **thread\_function**: name of the function the thread should execute.
- **thread\_args**: where are the arguments to pass to thread function when it starts
  - Cannot share arguments between threads (data race, one thread modifies while others read).
- pthread\_create() **returns 0 if succeed**, and a nonzero error code otherwise.

```
for (i = 0; i < nthreads; i++) {  
    thread_ids[i] = i;  
    pthread_create(&thread_array[i], NULL, HelloWorld, &thread_ids[i]);  
}
```

# Joining thread

```
int pthread_join (pthread_t thread, void **retval);
```

- `pthread_join()` function suspends the execution of its caller until the thread it references terminates.
  - Called by the main thread to wait for the finish of worker threads.
  - When the main thread terminates, all other threads are terminated too.
- **thread**: which thread to wait on
- **retval**: where the thread's return value should be stored
  - `*retval` is a pointer returned by the thread function.

```
for (i = 0; i < nthreads; i++)  
    pthread_join(thread_array[i], NULL);
```

# Compile and run the program

```
$ gcc -o hellothreads hellothreads.c -pthread
```

```
$ ./hellothreads
```

```
usage: ./hellothreads <n>
```

```
where <n> is the number of threads
```

```
$ ./hellothreads 4
```

```
Hello world! I am thread 1
```

```
Hello world! I am thread 2
```

```
Hello world! I am thread 3
```

```
Hello world! I am thread 0
```

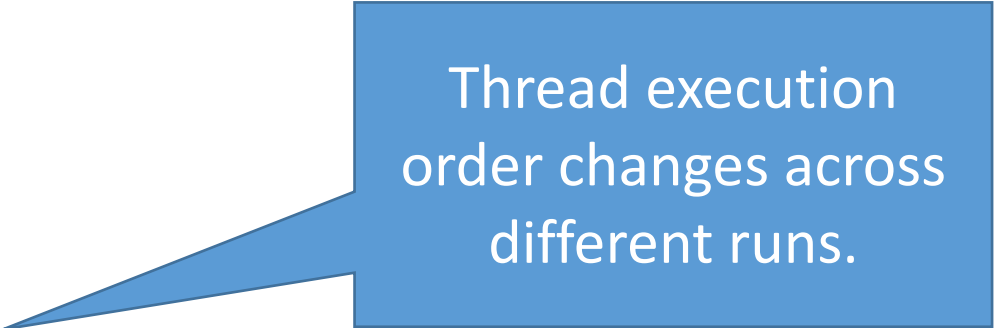
```
$ ./hellothreads 4
```

```
Hello world! I am thread 0
```

```
Hello world! I am thread 1
```

```
Hello world! I am thread 2
```

```
Hello world! I am thread 3
```



Thread execution  
order changes across  
different runs.




# Another example: scalar multiplication

## Sequential code of multiplying an array and an integer

```
void scalar_multiply(int * array, long length, int s) {  
    for (i = 0; i < length; i++)  
        array[i] = array[i] * s;  
}
```

To create a multi-threaded version of this application with  $t$  threads, it is necessary to:

- Create  $t$  threads.
- Assign each thread a subset of the input array (i.e.  $length/t$  elements).
- Instruct each thread to multiply the elements in its array subset by  $s$ .



```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

long *array, length, nthreads, s;

void *scalar_multiply(void *id) {
    long *myid = (long *) id;
    int i;

    //assign each thread its own chunk of elements to process
    long chunk = length / nthreads;
    long start = (*myid) * chunk;
    long end = start + chunk;

    if ( (*myid) == nthreads - 1) end = length;

    //perform scalar multiplication on assigned chunk
    for (i = start; i < end; i++) {
        array[i] *= s;
    }

    return NULL;
}
```

```

int main(int argc, char **argv) {
    int i;
    pthread_t *thread_array; //pointer to future thread array
    long *thread_ids;

    nthreads = strtol(argv[1], NULL, 10);
    length = strtol(argv[2], NULL, 10);
    s = strtol(argv[3], NULL, 10);
    array = (long *)malloc(sizeof(long)*length);
    thread_array = malloc(nthreads * sizeof(pthread_t));
    thread_ids = malloc(nthreads * sizeof(long));
    ... /* initialize array */
    for (i = 0; i < nthreads; i++) {
        thread_ids[i] = i;
        pthread_create(&thread_array[i], NULL,
                      scalar_multiply, &thread_ids[i]);
    }

    for (i = 0; i < nthreads; i++) pthread_join(thread_array[i], NULL);
    ... /* print array or save it into file */
    free(thread_array); free(thread_ids); free(array); return 0;
}

```

# Private data and shared data

- Global variables are shared by all threads.
- Local variables declared in thread function (stack variables) are private to each thread.
- Data shared between main thread and a worker thread
  - Input arguments of thread function passed via a pointer.
  - Data returned by thread function passed via another pointer.

# Avoid using global variables by passing data to threads

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

struct t_arg {
    int *array; // pointer to shared array
    long length; // num elements in array
    long s; //scaling factor
    long numthreads; // total number of threads
    long id; // logical thread id
};

void * scalar_multiply(void* args) {
    //cast to a struct t_arg from void*
    struct t_arg * myargs = (struct t_arg *) args;

    //extract all variables from struct
    long myid = myargs->id;
    long length = myargs->length;
    long s = myargs->s;
    long nthreads = myargs->numthreads;
    int * ap = myargs->array; //pointer to array in main
```



```
//code as before
```

```
    long chunk = length/nthreads;  
    long start = myid * chunk;  
    long end   = start + chunk;  
    if (myid == nthreads-1) end = length;  
  
    for (int i = start; i < end; i++) ap[i] *= s;  
    return NULL;  
}
```

```
int main(int argc, char **argv) {  
    int i;  
    pthread_t *thread_array; //pointer to future thread array  
    long *thread_ids;  
    long nthreads = strtol(argv[1], NULL, 10); //get number of threads  
    long length = strtol(argv[2], NULL, 10); //get length of array  
    long s = strtol(argv[3], NULL, 10); //get scaling factor  
    int *array = malloc(length*sizeof(long));  
    pthread_t *thread_array = malloc(nthreads * sizeof(pthread_t));  
    struct t_arg *thread_args = malloc(nthreads * sizeof(struct t_arg));  
}
```

**//Populate thread arguments for all the threads**

```
for (i = 0; i < nthreads; i++) {  
    thread_args[i].array = array;  
    thread_args[i].length = length;  
    thread_args[i].s = s;  
    thread_args[i].numthreads = nthreads;  
    thread_args[i].id = i;  
}
```

```
... /* initialize array */  
for (i = 0; i < nthreads; i++)  
    pthread_create(&thread_array[i], NULL,  
                  scalar_multiply, &thread_args[i]);
```

```
for (i = 0; i < nthreads; i++) pthread_join(thread_array[i], NULL);  
... /* print array or save it into file */  
free(thread_array); free(thread_ids); free(array); return 0;  
}
```

# common “bugs” that first-time pthread programmers make

- When creating multiple threads, pthread\_t struct and/or thread argument are reused across threads and overwritten.
- pthread\_join() is not called in main thread
  - the main thread may reach the end of main() and exit
  - other threads are terminated when main thread finishes
    - It possible that they have NOT had a chance to compute a result
- The following code is sequential!

```
for (i=0; i < num_threads; i++) {  
    pthread_create(&(threads[i]), ...)   
    pthread_join(threads[i], ...)   
}
```

# Synchronizing Threads

# Why is synchronization important?

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <assert.h>
#define NUM_THREADS 2
#define ITERATIONS_PER_THREAD 5000000
int cnt = 0;
void * worker( void *ptr )
{
    int i;
    for (i = 0; i < ITERATIONS_PER_THREAD; i++)
        cnt++;
}

int main(void)
{
    pthread_t threads[NUM_THREADS];
    int i, result;
    /* Start threads */
```

```

for (i = 0; i < NUM_THREADS; i++) {
    result = pthread_create(&threads[i], NULL,
                           worker, NULL);

    assert(result == 0);
}
/* Wait for threads to finish */
for (i = 0; i < NUM_THREADS; i++) {
    result = pthread_join(threads[i], NULL);
    assert(result == 0);
}
printf("Final value: %d (%.2f%%)\n", cnt,
       100.0 * cnt / (NUM_THREADS *
                     (double)ITERATIONS_PER_THREAD));
}

```

What will be printed out if everything works as expected?

What will be printed out actually?

How small/large can the final value be?

# Shared memory “code” for increasing a counter

```
int cnt = 0;
```

Thread 1

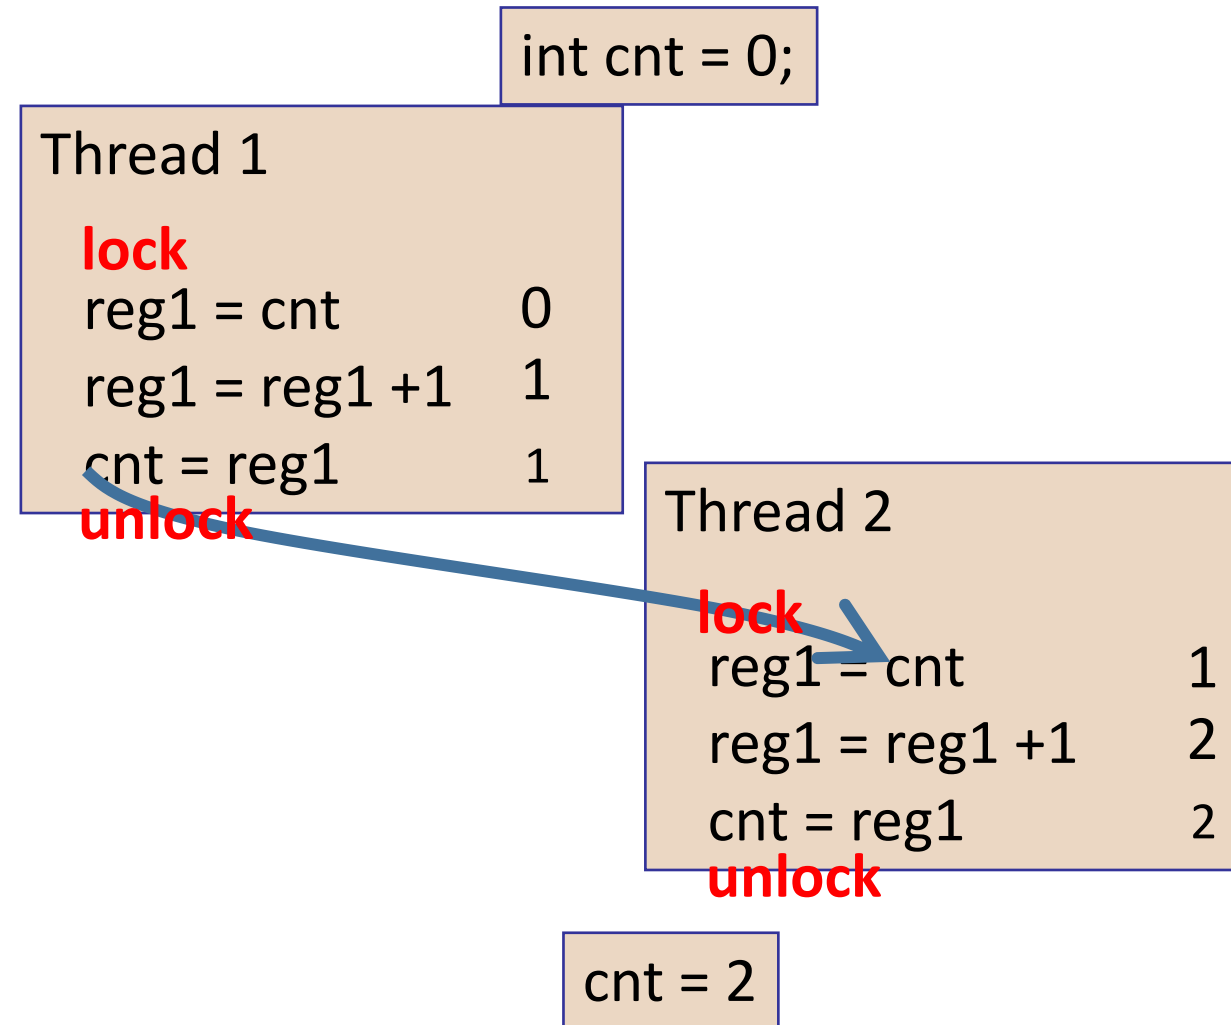
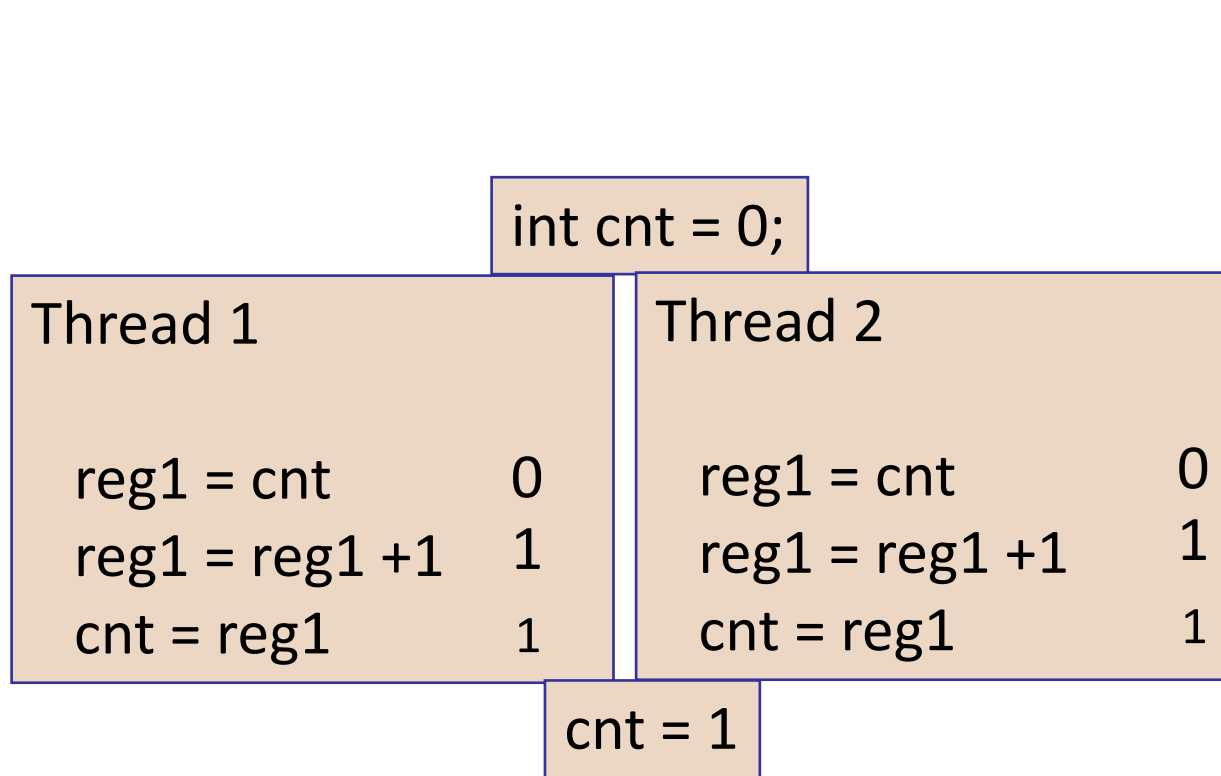
```
for (i = 0, i<100; i++)  
    cnt++;
```

Thread 2

```
for (i = 0, i<100; i++)  
    cnt++;
```

- Problem is a race condition on the shared variable ***cnt*** in the program
- A race condition or data race occurs when:
  - two threads share the same variable, and at least one does a write.
  - The accesses are not synchronized so they could happen simultaneously

# Shared memory “code” for increasing a counter





# Semaphores

- A non-negative global integer synchronization variable
- Manipulated by wait and post operations:
  - `wait(s): [ while (s == 0) wait(); s--; ]`
    - Also `P(s)`, Dutch for "Proberen" (test)
  - `post(s): [ s++; ]`
    - Also `V(s)`, Dutch for "Verhogen" (increment)
- OS kernel guarantees that operations between brackets `[ ]` are executed indivisibly
  - i.e., `s--` can't be broken into load/update/store
  - Result: only one wait or post operation at a time can modify `s`
  - When while-loop in wait terminates, only that wait can decrement `s`
  - `S` never goes below 0

# Semaphores (unnamed)

- Include header file `semaphore.h`
- Declare semaphore as global variable and thus shared between threads
  - `sem_t mysem;`
- Initialize the semaphore using **`sem_init()`** (usually in `main()` ).
  - `int sem_init(sem_t *sem, int pshare, unsigned int value);`
  - `sem`: pointer to semaphore
  - `pshare`: 0 --- shared between threads; 1 --- shared between process
  - `value`: initial value of semaphore
- Destroyed using **`sem_destroy()`**
  - `int sem_destroy(sem_t *sem);`
- Wait operation using **`sem_wait()`**, and Post operation using **`sem_post()`**
  - `int sem_wait(sem_t *s);`
  - `int sem_post(sem_t *s);`

```
#include <semaphore.h>

...
int cnt = 0;
sem_t cnt_sem;
void * worker( void *ptr ){
    int i;
    for (i = 0; i < ITERATIONS_PER_THREAD; i++) {
        sem_wait(&cnt_sem);
        cnt++;
        sem_post(&cnt_sem);
    }
}

int main(void){
    ...
    result = sem_init(&cnt_sem, 0, 1); /* Initialize semaphore */
    if (result < 0)
        exit(-1);
    ...
    sem_destroy(&cnt_sem);
}
```

# Mutexes in pthreads

- If possible, use mutexes over semaphores
  - Mutexes: only locking thread can unlock
  - Semaphores: any thread can decrement (unlock); much harder to manage
- To declare a mutex: `pthread_mutex_t amutex;`
- To initialize a mutex:  
`amutex = PTHREAD_MUTEX_INITIALIZER;` or  
`pthread_mutex_init(&amutex, NULL);`
- To use it:  
`int pthread_mutex_lock(&amutex);`  
`int pthread_mutex_unlock(&amutex);`
- To destroy a mutex  
`int pthread_mutex_destroy(&amutex);`

# Barrier -- global synchronization

- A point in a program where all threads must reach before any thread can cross
  - threads reach the barrier & then wait until all other threads arrive
  - all threads reach & begin executing code beyond the barrier
- simple use of barriers -- all threads hit the same one

- ```
for(...){  
    work_on_my_problem() ;  
    barrier;  
    get_data_from_others() ;  
    barrier;  
}
```

- more complicated -- barriers on branches (or loops)

```
if (tid % 2 == 0) {  
    work1() ;  
    barrier  
}  
else {  
    barrier  
}
```

# Creating and Initializing a Barrier

- To declare a barrier:

```
pthread_barrier_t b;
```

- To initialize a barrier for 3 threads:

```
pthread_barrier_init(&b, NULL, 3); or  
b=PTHREAD_BARRIER_INITIALIZER(3);
```

- To wait at a barrier:

```
pthread_barrier_wait(&b);
```

- To destroy a barrier

```
pthread_barrier_destroy(&b);
```

A more complicated example  
(multi-threaded CountSort)

# CountSort

- Designed for sorting many integer values within a small range.
- For an array, CountSort counts the frequency of each value in the array, and then enumerates each value by its frequency.

```
void countElems(int *counts,  
int *array_A, long length) {  
  
    int val, i;  
    for(i = 0; i < length; i++) {  
        val=array_A[i];  
        counts[val]=counts[val]+1;  
    }  
}
```

A = [9, 0, 2, 7, 9, 0, 1, 4, 2, 2, 4, 5, 0, 9, 1]




Count the frequency of each  
value from 0 to 9

counts = [3, 2, 3, 0, 2, 1, 0, 1, 0, 3]

2 1s

3 0s



Emulate each value by  
its frequency

A = [0, 0, 0, 1, 1, 2, 2, 2, 4, 4, 5, 7, 9, 9, 9]



# How to parallelize countElems?

```
pthread_mutex_t mutex;  
void *countElems( void *args ) {
```

```
    struct t_arg * myargs = (struct t_arg *)args;
```

```
    int *array = myargs->ap;
```

```
    long *counts = myargs->countp;
```

```
    long chunk = length / nthreads, start = myid * chunk;
```

```
    long end = (myid + 1) * chunk, val, i;
```

```
    if (myid == nthreads - 1) end = length;
```

```
    pthread_mutex_lock(&mutex);
```

```
    for (i = start; i < end; i++) {
```

```
        val = array[i];
```

```
        counts[val] = counts[val] + 1;
```

```
    }
```

```
    pthread_mutex_unlock(&mutex);
```

```
    return NULL;
```

```
}
```

[https://diveintosystems.org/singlepage/book/modules/SharedMemory/assets/attachments/countElems\\_p\\_v2.c](https://diveintosystems.org/singlepage/book/modules/SharedMemory/assets/attachments/countElems_p_v2.c)

[https://diveintosystems.org/singlepage/book/modules/SharedMemory/assets/attachments/countElems\\_p\\_v3.c](https://diveintosystems.org/singlepage/book/modules/SharedMemory/assets/attachments/countElems_p_v3.c)

**Computation is still serial.  
No speed up.**

Each thread accumulate frequencies into a local array  
and then accumulate local frequencies to global array

[https://diveintosystems.org/singlepage/book/modules/SharedMemory/assets/attachments/countElems\\_p\\_v3.c](https://diveintosystems.org/singlepage/book/modules/SharedMemory/assets/attachments/countElems_p_v3.c)

```
void *countElems( void *args ) {  
    struct t_arg * myargs = (struct t_arg *)args;  
    int *array = myargs->ap, *counts = myargs->countp;  
    long local_counts[MAX] = {0};  
  
    long chunk = length / nthreads, start = myid * chunk;  
    long end = (myid + 1) * chunk, val, i;  
    if (myid == nthreads-1) end = length;  
  
    for (i = start; i < end; i++) {  
        val = array[i];  
        local_counts[val] = local_counts[val] + 1;  
    }  
    pthread_mutex_lock(&mutex);  
    for (i = 0; i < MAX; i++) counts[i] += local_counts[i];  
    pthread_mutex_unlock(&mutex);  
    return NULL;  
}
```


Most  
computation  
can be done  
in parallel.

# CS 288 Intensive Programming in Linux


Professor Ding, Xiaoning

# Two architectures and two programming models

## Shared Memory System



## Distributed Memory System



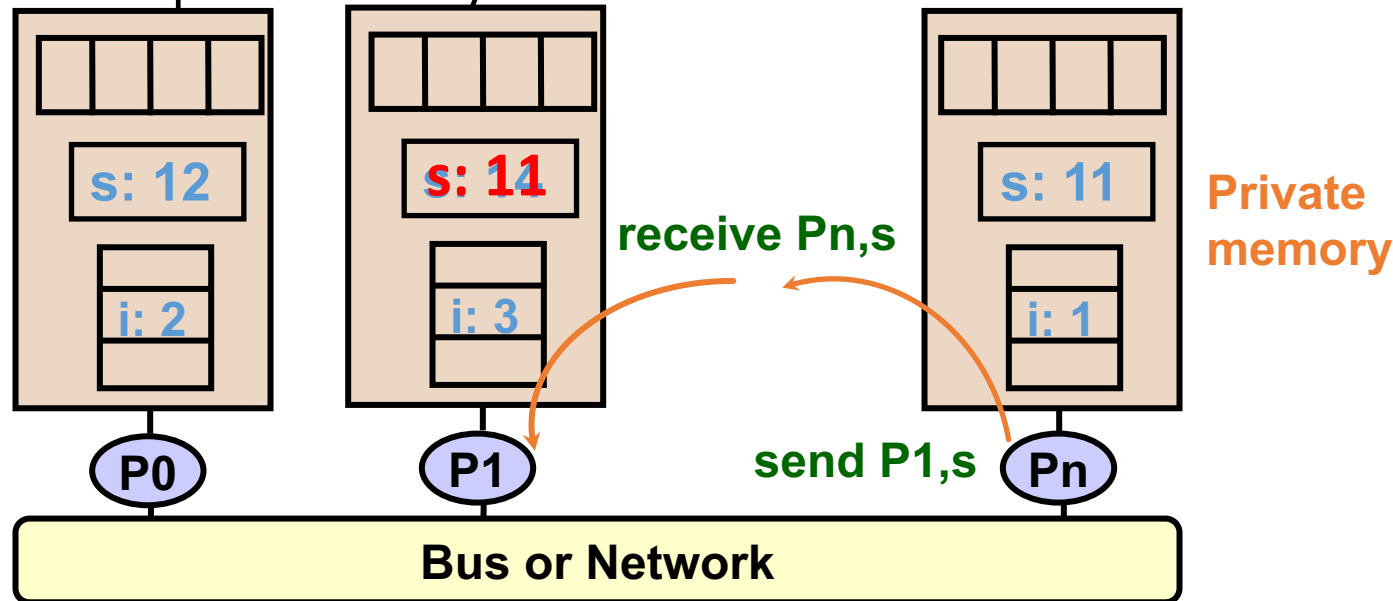
- Shared memory programming
  - pthread, open MP
  - Usually used on shared memory systems
- Message passing programming
  - processes, open MPI
  - Can be used on both types of systems
- Using multiple computers to increase computing capacity and speed
  - More sources: CPU cores, memory capacity, storage bandwidth, etc.
  - Many scenarios need super computing, e.g., fluid dynamics simulation, AI training, etc.

# What is Message Passing Interface(MPI)

- A standard interface for implementing message-passing libraries:
  - Extended message-passing model
  - Not a language or compiler specification, not a specific implementation or product
  - 125 functions, but only 6 basic functions.
  - Applications can be written in C, C++ or Fortran and appropriate calls to MPI can be added where required
- Software packages implementing the interface and the related tools.
  - Usually open-source, mainly includes library for message passing, compiler wrapper, and supporting tools.
  - Dominant MPI implementations on Linux: Open MPI and MPICH
  - Other implementations: Microsoft MPI
- Hide the complexity and difference of parallel computing architectures
  - Develop one program, run on different systems --- super computers, cluster of heterogeneous servers, etc.

# MPI program

- A program consists of a collection of processes.
  - Usually created at program startup time and kept fixed during execution
- SPMD (Single Program Multi-Data)
  - Multiple processes execute the same copy of code
  - Each process is a thread of control with its own local address space.
- NO shared data between processes.
  - Processes communicate by explicit send and receive function pairs
  - Coordination is implicit in every communication event.



# Basic MPI functions and commands

# Simple MPI program identifying processes

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char ** argv)
{
    int rank, size;

    MPI_Init(&argc, &argv);

    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    printf("I am %d of %d\n", rank + 1, size);

    MPI_Finalize();
    return 0;
}
```

*Basic  
requirements  
for an MPI  
program*



# Basic requirements for an MPI program

- Include header file `mpi.h`.
- `MPI_Init(&argc, &argv)` starts MPI and a process group  
--- `MPI_COMM_WORLD`
- `MPI_Finalize()` exits MPI
- The group of processes in the program is called `MPI_COMM_WORLD`
  - `MPI_Comm_size(MPI_COMM_WORLD, &size)` gets the size of the group, i.e., # of processes.
- Each process is identified by its *rank*.
  - `MPI_Comm_rank(MPI_COMM_WORLD, &rank)` gets the rank of the caller process.

# Compile and Run a MPI program

- **Compilation:**

- Regular applications: `gcc test.c -o test`
- MPI applications: `mpicc test.c -o test`

- **Execution:**

- Regular applications: `./test`
- MPI applications (running with N processes):

```
$mpirun -n 4 ./test
```

```
I am 1 of 4
```

```
I am 2 of 4
```

```
I am 3 of 4
```

```
I am 4 of 4
```

# Run a MPI program on multiple computers

- The program must be accessible by all computers with the same pathname.
  - Use a NFS or manually copy the file (scp).
- The user can remotely run a program on these computers using ssh
  - Better to set up password-less ssh login on all computers.

- Run 16 processes on 4 computers (4 processes on each of h1 ~ h3)

```
$mpirun -hosts h1:4,h2:4,h3:4,h4:4 -n 16 ./test
```

- Runs the first four processes on h1, the next four on h2, etc.

```
$mpirun -hosts h1,h2,h3,h4 -n 16 ./test
```

- Runs the first process on h1, the second on h2, etc., and wraps around
- So, h1 will have the 1<sup>st</sup>, 5<sup>th</sup>, 9<sup>th</sup> and 13<sup>th</sup> processes

- If there are many nodes, it might be easier to create a host file

```
$cat hf
```


```
h1:14
```

```
h2:2
```

```
$mpirun -hostfile hf ./test
```

# MPI program

- A program consists of a collection of processes.
  - Usually created at program startup time and kept fixed during execution
- SPMD (Single Program Multi-Data)
  - Multiple processes execute the same copy of code
  - Each process is a thread of control with its own local address space.
- NO shared data between processes.
  - Processes communicate by explicit send and receive function pairs
  - Coordination is implicit in every communication event.



# Single Program Multi-Data (SPMD)

- Multiple processes share the same copy of code.
  - Program is actually launched separately on different computers.
- Each process is a thread of control with its own local address space.
  - Each process has a copy of `rank`.
  - `MPI_Comm_rank(MPI_COMM_WORLD, &rank)` saves different values to different `rank` copies in different processes.
  - `printf()` prints different `rank` values by different processes.
  - Though different processes have different copies of `size` variable, `MPI_Comm_size(MPI_COMM_WORLD, &size)` saves the same value into them.
- Different processes can execute different code based on their ranks.
- Different processes can exchange data using messages.
  - `MPI_Send()`
  - `MPI_Recv()`
  - ...

# Different processes can execute different code and exchange data

*if (rank == 0) { // code executed by process 0*

prepare data

**MPI\_Send**(..**data** to rank 1..)

**MPI\_Recv**(..**results** from rank 1..)

*else if (rank == 1) { //code for process 1*

**MPI\_Recv**(..**data** from rank 0..)

compute results

**MPI\_Send**(..**results** to rank 0..)

}

- MPI\_Send() sends a message
- MPI\_Recv() waits and receives a message
- MPI\_Send() and MPI\_Recv() must be properly paired for smooth execution

Offload some computation from process 0 to process 1

**Process 0 (rank 0):**

Prepare **data**

**MPI\_send**(...**data** to rank **1**...)

**MPI\_recv**(...**results** from rank **1**...)

**Process 1 (rank 1)**

**MPI\_recv**(...**data** from rank **0**...)


Calculate **results** from data

**MPI\_send**(...**results** to rank **0**...)

# Matrix multiplication in parallel


```

all computation
for (i=0;i<N;i++)
  for (j=0;j<N;j++)
    for (k=0;k<N;k++)
      C[i][j] += A[i][k]
                *B[k][j]
    
```



```

computation of each worker
for (i=low;i<high;i++)
  for (j=0;j<N;j++)
    for (k=0;k<N;k++)
      C[i][j] += A[i][k]
                *B[k][j]
    
```



# Matrix multiplication in parallel (1 master multiple workers)

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
#define N 100          /* size of each dimension */
#define FROM_MASTER 1  /* setting a message type */
#define FROM_WORKER 2  /* setting a message type */

int main (int argc, char *argv[]) {
    int    numtasks, taskid, numworkers, rows, i, j, k;
    double a[N][N], b[N][N], c[N][N]; /*matrices */
    MPI_Status status;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &taskid);
    MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
    numworkers = numtasks - 1;
    rows = N/numworkers;
```



```

if (taskid == 0)    { /* code for master */
    /*Initializing arrays*/
    for (i=0; i<N; i++)
        for (j=0; j<N; j++) {    a[i][j]= i+j;    b[i][j]= i*j;    }

    /* Send matrix data to the worker tasks */
    for (int dest=1; dest<=numworkers; dest++)    {
        MPI_Send(&a[ (dest-1)*rows][0], rows*N,
                 MPI_DOUBLE, dest, FROM_MASTER, MPI_COMM_WORLD);
        MPI_Send(&b, N*N, MPI_DOUBLE, dest,
                 FROM_MASTER, MPI_COMM_WORLD);
    }
    /* Receive results from worker tasks */
    for (int source=1; source<=numworkers; source++)
        MPI_Recv(&c[ (source-1)*rows][0], rows*N, MPI_DOUBLE,
                 source, FROM_WORKER, MPI_COMM_WORLD, &status);

    printf("***** Result Matrix *****\n");
}

```

```

    for (i=0; i<N; i++) {
        for (j=0; j<N; j++) printf("%6.2f  ", c[i][j]);
        printf("\n");
    }
}
if (taskid > 0) { /* code for worker */
    MPI_Recv(&a, rows*N, MPI_DOUBLE, 0, FROM_MASTER,
             MPI_COMM_WORLD, &status);
    MPI_Recv(&b, N*N, MPI_DOUBLE, 0, FROM_MASTER,
             MPI_COMM_WORLD, &status);
    for (i=0; i<rows; i++)
        for (j=0; j<N; j++)
            for (k=0; k<N; k++)
                c[i][k] = c[i][k] + a[i][j] * b[j][k];

    MPI_Send(&c, rows*N, MPI_DOUBLE, 0, FROM_WORKER,
             MPI_COMM_WORLD);
}
MPI_Finalize();
}

```

MPI\_Send(void \*buf, int count, MPI\_Datatype datatype, int dest, int tag,  
MPI\_Comm comm)

MPI\_Recv(void \*buf, int count, MPI\_Datatype datatype, int src, int tag,  
MPI\_Comm comm, MPI\_Status \*stat)

- buf, count, and datatype determine what data would be sent/received.
- Information for pairing sender and receiver: comm, dest, src
  - Use MPI\_COMM\_WORLD as comm.
  - For sender process, dest in MPI\_Send() is where should the data be sent to (i.e., receiver process).
  - For receiver process, src in MPI\_Recv() is where should it expect the data coming from (i.e., sender process).
- tag is a user-defined message identification number (integer) for pairing MPI\_Send() and MPI\_Recv() calls.
- When MPI\_Send() returns, the data has been delivered to the system and the buffer can be reused.
- When MPI\_Recv() returns, the data is received and saved in buf.

# Use MPI data types for communications

## **MPI datatype**

MPI\_CHAR

MPI\_SIGNED\_CHAR

MPI\_UNSIGNED\_CHAR

MPI\_SHORT

MPI\_UNSIGNED\_SHORT

MPI\_INT

MPI\_UNSIGNED

MPI\_LONG

MPI\_UNSIGNED\_LONG

MPI\_FLOAT

MPI\_DOUBLE

MPI\_LONG\_DOUBLE

## **C datatype**

signed char

signed char

unsigned char

signed short

unsigned short

signed int

unsigned int

signed long

unsigned long

float

double


long double

MPI datatype is very similar to a C

# MPI\_Send/MPI\_Recv can be replaced by collective operations

MPI\_Bcast(void \***buffer**, int **count**, MPI\_Datatype **type**, int **root**, MPI\_Comm **comm**)

- broadcasts message from root to all processes (including root).
- Called by all processes
- INOUT : buffer (starting address)
- IN : count (num entries in buffer)
- IN : type (type of each entry)
- IN : root (rank of broadcast root)
- IN : comm (communicator, MPI\_COMM\_WORLD)
- Improve both simplicity and efficiency.
- On return, contents of buffer is copied to all processes.



```
MPI_Reduce(void *sendbuf, void *recvbuf, int count,  
MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm)
```

- IN sendbuf (address of send buffer)
  - OUT recvbuf (address of receive buffer)
  - IN count (number of elements in send buffer)
  - IN datatype (data type of elements in send buffer)
  - IN op (reduce operation)
  - IN root (rank of root process)
  - IN comm (communicator)
- 
- MPI\_Reduce combines elements specified by send buffer and performs a **reduction operation** op on them.

# MPI built-in collective computation operations

- **MPI\_MAX** Maximum
- **MPI\_MIN** Minimum
- **MPI\_PROD** Product
- **MPI\_SUM** Sum
- **MPI\_LAND** Logical and
- **MPI\_LOR** Logical or
- **MPI\_LXOR** Logical exclusive or
- **MPI\_BAND** Bitwise and
- **MPI\_BOR** Bitwise or
- **MPI\_BXOR** Bitwise exclusive or
- **MPI\_MAXLOC** Maximum and location
- **MPI\_MINLOC** Minimum and location

# Example: Calculating the approximation of pi

```
#include "mpi.h"
#include <math.h>
#include <stdlib.h>
#include <stdio.h>
```

$$PI = \lim_{n \rightarrow \infty} \left( \frac{1}{n} \sum_{i=1}^n \frac{4.0}{1.0 + \left( \frac{i-0.5}{n} \right)^2} \right)$$

```
int main(int argc, char *argv[])
{
    int n, myid, numprocs, i;
    double PI25DT = 3.141592653589793238462643;
    double pi, sum, partial_sum, x;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);
```



```

if (myid == 0) {
    printf("Enter the number of terms:");
    scanf("%d",&n);
}
MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);

partial_sum = 0.0;
for(i=n/numprocs*myid+1; i<= n/numprocs*(myid+1); i++) {
    x = ((double)i - 0.5)/n;
    partial_sum+= 4.0 / (1.0 + x*x);
}


MPI_Reduce(&partial_sum, &sum, 1, MPI_DOUBLE, MPI_SUM, 0,
          MPI_COMM_WORLD);

if (myid == 0){
    pi = sum/n;
    printf("pi is approximately %.16f, Error is %.16f\n",
        pi, fabs(pi - PI25DT));
}
MPI_Finalize();
return 0;
}

```

# MPI\_Gather and MPI\_Scatter

- MPI\_Scatter() distribute data from one process to all other processes in a communicator
  - Inverse of MPI\_Gather
  - Data elements on root listed in rank order – each processor gets corresponding data chunk after call to scatter.
- MPI\_Gather() gathers together values from a group of processes
  - Each process sends content of send buffer to the root process.
  - Root receives and stores in rank order.



```
MPI_Scatter(void *sendbuf,int sendcount,MPI_Datatype sendtype,  
           void *recvbuf, int recvcount, MPI_Datatype recvtype,  
           int root, MPI_Comm comm)
```

- IN sendbuf (starting address of send buffer)
- IN sendcount (number of elements **sent to each process**)
- IN sendtype (type)
- OUT recvbuf (address of receive bufer)
- IN recvcount (n-elements **in receive buffer**)
- IN recvtype (data type of receive elements)
- IN root (rank of sending process)
- IN comm (communicator)

**Note:** **send\_buf, sendcount, sendtype** arguments ignored for non-root processes.

```
MPI_Gather(void *sendbuf,int sendcount,MPI_Datatype sendtype,  
          void *recvbuf,int recvcount,MPI_Datatype recvtype,  
          int root, MPI_Comm comm)
```

- IN sendbuf (starting address of send buffer)
- IN sendcount (number of elements **in send buffer**)
- IN sendtype (type)
- OUT recvbuf (address of receive bufer)
- IN recvcount (n-elements **for any single receive**)
- IN recvtype (data type of recv buffer elements)
- IN root (rank of receiving process)
- IN comm (communicator, MPI\_COMM\_WORLD)

**Note**: **Recvbuf** argument ignored for all non-root processes (also recvtype, etc.)  
**recvcount** on root indicates number of items received from each process, not total. This is a very common error.

# Matrix multiplication in parallel (each process works on one data piece)

```
#include "mpi.h"
#include <stdio.h>
#include <stdlib.h>

#define N 100                /* size of each dimension */

int main (int argc, char *argv[]) {
    int numtasks, taskid, rows, i, j, k;
    double a[N][N], b[N][N], c[N][N], aprime[N][N], cprime[N][N];
    MPI_Status status;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &taskid);
    MPI_Comm_size(MPI_COMM_WORLD, &numtasks);

    if (taskid == 0) {
        printf("run with %d tasks.\n", numtasks);
        for (i=0; i<N; i++) /*Initializing arrays*/
            for (j=0; j<N; j++) { a[i][j]= i+j; b[i][j]= i*j; }
    }
}
```

```

rows = N/numtasks;
MPI_Bcast(b, N*N, MPI_DOUBLE, 0, MPI_COMM_WORLD);
MPI_Scatter(a, rows * N, MPI_DOUBLE, aprime, rows * N, MPI_DOUBLE,
            0, MPI_COMM_WORLD);

for (i=0; i<rows; i++)
    for (k=0; k<N; k++)
        for (j=0; j<N; j++)
            cprime[i][k] = cprime[i][k] + aprime[i][j] * b[j][k];

MPI_Gather(cpime, rows * N, MPI_DOUBLE, c, rows * N, MPI_DOUBLE,
            0, MPI_COMM_WORLD);

if(taskid == 0) {
    printf("***** Result *****\n");
    for (i=0; i<N; i++) {
        for (j=0; j<N; j++) printf("%6.2f  ", c[i][j]);
        printf("\n");
    }
}
MPI_Finalize();
}

```

# Avoid deadlock

- Process may wait in MPI\_Send() and MPI\_Recv().
- Deadlock may be caused if they are not placed in a correct order.

## Process 0 (rank 0):

**MPI\_recv**(...**B** from rank **1**...)

**MPI\_send**(...**A** to rank **1**...)

## Process 1 (rank 1)

**MPI\_recv**(...**A** from rank **0**...)

**MPI\_send**(...**B** to rank **0**...)

# CS 288 Intensive Programming in Linux

Professor Ding, Xiaoning



# Data in a file is also a bit string (no type info)

Check file content using hexadecimal format

Vim:%!xxd to switch to hex, :%!xxd -r to switch back

| address area | hexadecimal area |      |      |      |      |      |      |      |      |  | character area   |
|--------------|------------------|------|------|------|------|------|------|------|------|--|------------------|
| 00000000:    | 7f45             | 4c46 | 0201 | 0100 | 0000 | 0000 | 0000 | 0000 | 0000 |  | .ELF.....        |
| 00000010:    | 0300             | 3e00 | 0100 | 0000 | 5058 | 0000 | 0000 | 0000 | 0000 |  | ..>.....PX.....  |
| 00000020:    | 4000             | 0000 | 0000 | 0000 | a003 | 0200 | 0000 | 0000 | 0000 |  | @.....           |
| 00000030:    | 0000             | 0000 | 4000 | 3800 | 0900 | 4000 | 1c00 | 1b00 |      |  | ....@.8...@..... |
| 00000040:    | 0600             | 0000 | 0500 | 0000 | 4000 | 0000 | 0000 | 0000 |      |  | .....@.....      |
| 00000050:    | 4000             | 0000 | 0000 | 0000 | 4000 | 0000 | 0000 | 0000 |      |  | @.....@.....     |
| 00000060:    | f801             | 0000 | 0000 | 0000 | f801 | 0000 | 0000 | 0000 |      |  | .....            |

**Data contents in /bin/l**

# vi /usr/share/doc/gcc/README.Debian

What does it look like normally:

```
The Debian GNU Compiler Collection Setup
=====
```

Abstract

-----

Debian uses a default version of GCC for most packages; however, some

What is saved in the file:

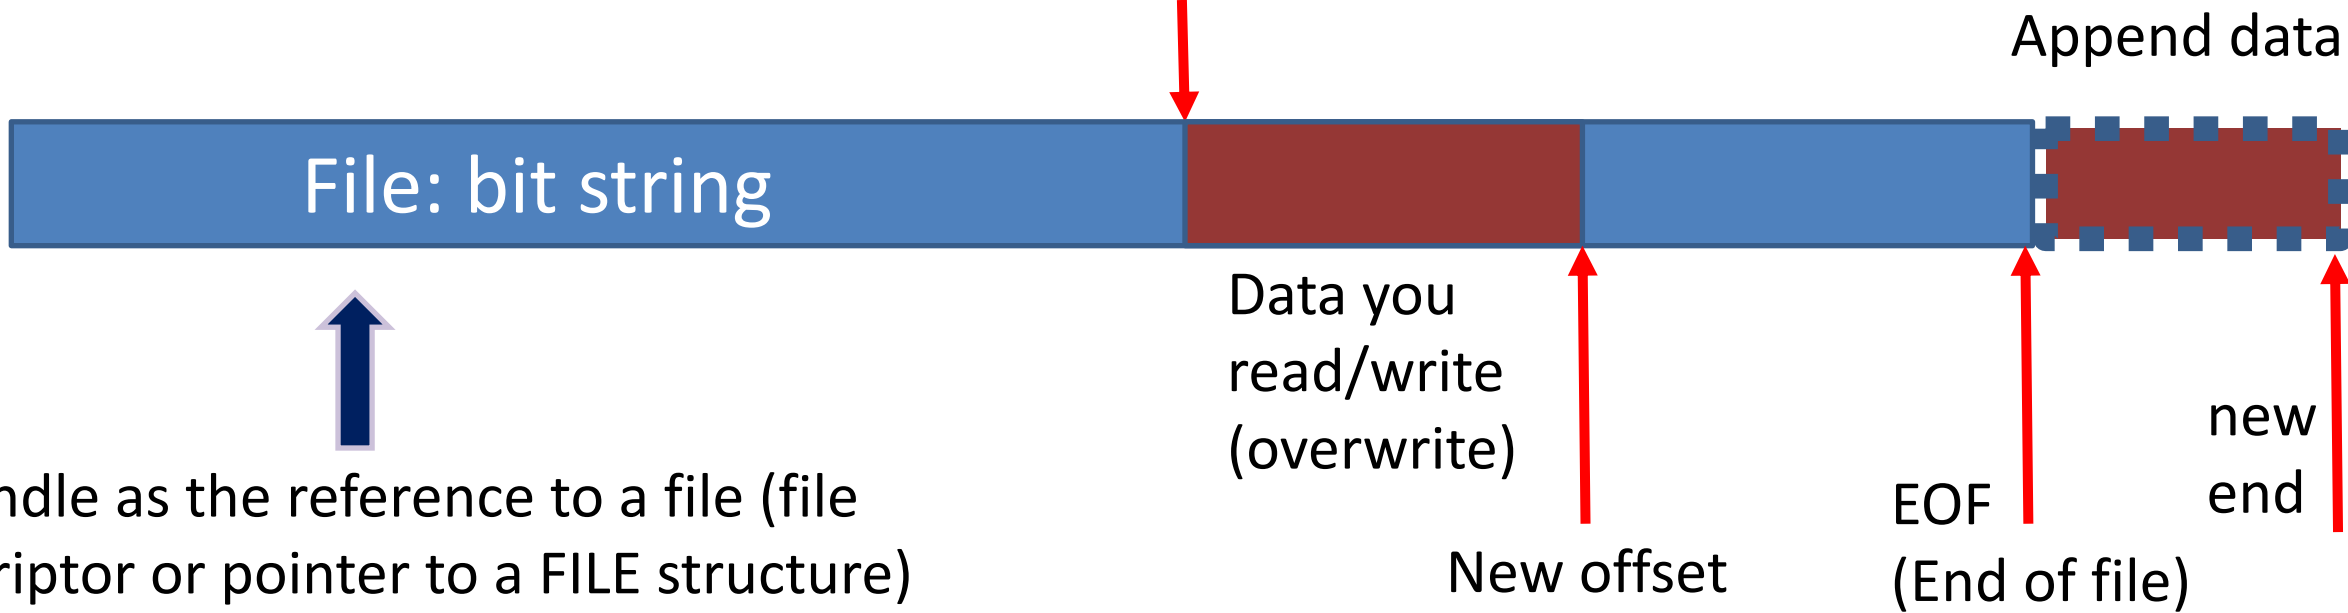
```
00000000: 0909 5468 6520 4465 6269 616e 2047 4e55  ..The Debian GNU
00000010: 2043 6f6d 7069 6c65 7220 436f 6c6c 6563  Compiler Collec
00000020: 7469 6f6e 2053 6574 7570 0a09 093d 3d3d  tion Setup...==
00000030: 3d3d 3d3d 3d3d 3d3d 3d3d 3d3d 3d3d 3d3d  =====
```

# Outline

- File operations
- handle directory entries
- Travels a directory

# Overview of files and file operations

offset (where you can read/write)



A handle as the reference to a file (file descriptor or pointer to a FILE structure)

## Step 1: Open a file

- Information needed: pathname and mode
- Can be used to create a file
  - Additional info needed: file permissions.
- Return a handle

## Step 3: close a file

## Step 2: Read/write a file

- Happens at the current file offset
  - Change file offset when necessary
- When offset is in the middle of the file
  - returns the number of bytes
  - Overwrites data
  - offset updated automatically
- When offset is at the end of a file, reads return EOF and writes append

# Library functions for opening, creating, and closing a file

```
#include <stdio.h>
FILE * fopen (const char *path, const char *mode);
int fclose(FILE *stream);
```

```
FILE *stream;
stream = fopen ("/etc/manifest", "r");
if (stream == NULL)
    perror("fopen");
/* read data */
fclose(stream);
```

- return a pointer to FILE structure (FILE \*)
  - return NULL if file is not opened.
- If used to create a file, file permission is 666.

# Mode is a string

- r Open the file for reading.
  - The file offset is initialized to 0.
- r+ Open the file for both reading and writing.
  - The file offset is initialized to 0.
- w Open the file for writing.
  - If the file exists, it is truncated to zero length.
  - If the file does not exist, it is created.
  - The file offset is initialized to 0
- w+ Open the file for both reading and writing.
  - If the file exists, it is truncated to zero length.
  - If the file does not exist, it is created.
  - The file offset is initialized to 0
- a Open the file for appending.
  - If the file does not exist, it is created.
  - The file offset is initialized to the end of the file
- a+ Open the file for reading and appending.
  - If the file does not exist, it is created.
  - The file offset for reading is initialized to 0
  - Output is always written at the end of the file

# Library functions for changing file offset

```
#include <stdio.h>
int fseek (FILE *stream, long pos, int origin);
/* to get the current offset */
long ftell (FILE *stream);
/* to reset current offset to beginning of file */
void rewind(FILE *stream);
```

- the `origin` argument can be one of the following:
  - `SEEK_CUR`: new file offset = current value + pos
    - When `pos == zero`, `fseek` returns the current file offset.
  - `SEEK_END`: new file offset = the current length of the file + pos.
  - `SEEK_SET`: new file offset = pos.
- `pos` can be positive, zero, or negative.

# Library functions for reading/writing a file

```
#include <stdio.h>

/* read/write raw data */
size_t fread (void *buf, size_t size, size_t nr, FILE *stream);
size_t fwrite (void *buf, size_t size, size_t nr, FILE *stream);
/* read/write chars/strings */
int getc(FILE *stream);
int fgetc(FILE *stream);
char *fgets(char *s, int size, FILE *stream);
int fputc (char c, FILE *stream);
int fputs (const char *str, FILE *stream);
/* Read and convert into text */
int fscanf(FILE *stream, const char *format, ...);
/* convert into text and then write */
int fprintf(FILE *stream, const char *format, ...);
```



# A few questions on binary file vs. text file

- How can you tell whether a file is a binary file or text file?
- When you create a file, which format do you choose, binary or text?
- When you read/write a binary file, which functions do you choose?
- When you read/write a text file, which functions do you choose?

# Outline

- File operations
- Handle directory entries
- Travels a directory

# File and directory information

```
ubuntu@ip-172-30-0-5:~/temp$ ls -l
total 232
drwxrwx--- 2 ubuntu ubuntu 4096 Jan 24 05:13 bak
-rw--w---- 1 ubuntu ubuntu 116181 Nov 20 20:06 stock.html
-rw--w---- 1 ubuntu ubuntu 90722 Jul 8 2011 tagsoup-1.2.1.jar
-rwxrwx--- 1 ubuntu ubuntu 13728 Dec 15 19:38 test
-rw-rw---- 1 ubuntu ubuntu 2324 Dec 15 19:23 test.c
```

type

permissions

# of hard links

user & group  
names

size

last modification  
time

name

# Directory and directory entry

## Directory entry

```
struct dirent {  
    ino_t    d_ino;  
    /* inode number */  
    char    d_name[];  
    /* filename */  
    int len;  
    ...  
};
```

## Directory: a set of directory entries

```
struct dirent  
entry1={123, ".", ...}
```

```
struct dirent  
entry2={234, "..", ...}
```

```
struct dirent  
entry3={235, "photos", ...}
```

```
struct dirent  
entry4={236, "sort.c", ...}
```

# Methods for accessing directory entries

```
#include <sys/types.h>
```

```
#include <dirent.h>
```

```
DIR *opendir(const char *name);
```

```
struct dirent *readdir(DIR *dirp);
```

```
int closedir(DIR *dirp);
```

```
struct dirent {  
    ino_t    d_ino;  
    /* inode number */  
    char    d_name[256];  
    /* filename */  
    ...  
};
```

To get all the entries in a directory

- open a directory
- read out an entry (repeat until readdir returns NULL)
- Close the directory.

```

#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <sys/types.h>
#include <dirent.h>

int main(int argc, char **argv) {
    struct dirent *pent;    DIR *pdir;

    if(argc!=2 || strlen(argv[1])==0) {
        printf("%s dir\n", argv[0]);    exit(1);
    }

    pdir = opendir(argv[1]);
    if(pdir==NULL) {
        printf("open directory fails.\n");    exit(1);
    }

    while( (pent=readdir(pdir)) !=NULL )
        printf("%s\n", pent->d_name);

    closedir(pdir);
}

```

```

$ ./myls /bin
.
..
grep
loadkeys
btrfs-image
gendata
sync

```

# Obtaining file information other than names

```
#include <sys/stat.h>
```

```
int stat( const char* pathname, struct stat* buf )
```

struct stat has the information about a file

## NAME

st\_dev

st\_ino

st\_mode

st\_nlink

st\_uid

st\_gid

st\_size

st\_atime

st\_mtime

st\_ctime

## MEANING

device number

**file ID** -- inode number

**type and permission** flags

hard-link count

**user ID**

**group ID**

**file size**

**last access time**

**last modification time**

**last status-change time**

# Some predefined macros

- `st_mode` as argument and return true( a value of 1 ) for the following file types:

| MACRO | RETURNS TRUE FOR FILE TYPE |
|-------|----------------------------|
|-------|----------------------------|

|                       |                          |
|-----------------------|--------------------------|
| <code>S_IFDIR</code>  | directory                |
| <code>S_IFCHR</code>  | character special device |
| <code>S_IFBLK</code>  | block special device     |
| <code>S_IFREG</code>  | regular file             |
| <code>S_IFFIFO</code> | pipe                     |

- The time fields decoded using the standard C library `asctime()` and `localtime()` functions.



# Outline

- File operations
- Handle directory entries
- Travels a directory

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <dirent.h>
#include <time.h>

void traverse(char *pathName ) { /* traverse a directory */
    struct stat statBuf;
    mode_t mode;
    int result, charsRead;
    DIR *pDir;
    struct dirent *pEnt;
    char fileName[1024];

    result = stat(pathName, &statBuf); /* Obtain file status */
    if ( result == -1 ) { /* Status was not available */
        fprintf( stderr, "Cannot stat %s \n", pathName );
        return;
    }
}
```

```

mode = statBuf.st_mode;  /* Mode of file */

if (S_ISREG(mode)) /* Regular file */
    printf("%s: size %lu bytes, mod. time = %s",
           pathName, statBuf.st_size,
           asctime(localtime(&statBuf.st_mtime)));

else if ( S_ISDIR( mode ) ) { /* Directory */
    printf("Entering directory %s\n", pathName);
    pDir=opendir(pathName); /* Open for reading */
    while( (pEnt = readdir(pDir)) != NULL ) {
        if ( strcmp(pEnt->d_name, ".") != 0 &&
             strcmp( pEnt->d_name, "..") != 0) {
            sprintf( fileName, "%s/%s", pathName,
                    pEnt->d_name );
            traverse(fileName);
        }
    }
    closedir(pDir);
}
}
}

```

```
int main(int argc, char **argv) {  
    if(argc!=2 || strlen(argv[1])==0) {  
        printf("%s dir\n", argv[0]);    exit(1);  
    }  
  
    traverse(argv[1]);  
}
```

How to implement the following?

- Non-recursive DFS traversal
- Non-recursive BFS traversal
- Recursive BFS traversal

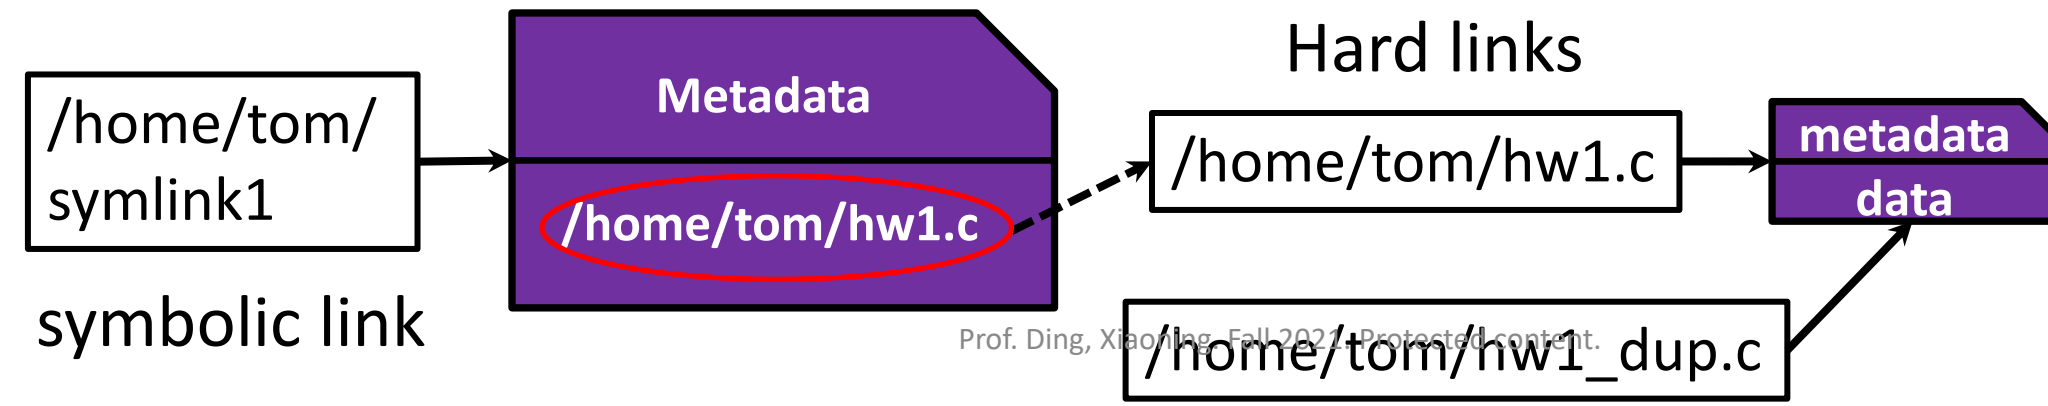

**ADDITIONAL MATERIALS ON FILE/DIRECTORY  
OPERATIONS (NOT REQUIRED, ONLY FOR INTERESTED  
STUDENTS)**

# Other methods

- directory entries
  - link
  - unlink
  - symlink
  - rename
- file metadata
  - chown
  - chmod

# What is a directory entry? Is it a file/subdirectory?

- A directory entry and the corresponding file
  - File → data and metadata (managing records)
  - Entry → pointer/link to the file.
- Two types of entries in a directory
  - Hard link (like a pointer to a file, or normal entries)
  - Symbolic link (symlink, like a pointer to hard link, or shortcuts in win systems)
- A file has only one and unique ID (**inode number**)
- A file may have multiple hard links and/or symbolic links



# Creating a new hard link

```
int link( const char* oldPath, const char*  
newPath)
```

- link() creates a new entry *newPath* and links it to the same file to which the label *oldPath* is linked.
  - The first hard link of a file is created when the file is created (i.e., pathname in open, create, or fopen call)
- The hard-link count of the associated file is incremented by one.
- If *oldPath* and *newPath* reside on different physical devices, a hard link cannot be made and link() fails.
- returns -1 if unsuccessful and a value of 0 otherwise



```

% cat mylink.c      ---> list the program.
main() {
    link("original.txt", "another.txt");
}

% cat original.txt  ---> list original file.
this is a file.

% ls -l original.txt another.txt  ---> another.txt not found
-rw-r--r--  1  glass           16  May 25 12:18  original.txt

% mylink

% ls -l original.txt another.txt  ---> examine files after.
-rw-r--r--  2  glass           16  May 25 12:18  another.txt
-rw-r--r--  2  glass           16  May 25 12:18  original.txt

% cat >> another.txt  ---> alter "another.txt".
hi
^D

% ls -l original.txt another.txt  ---> both labels reflect the change.
-rw-r--r--  2  glass           20  May 25 12:19  another.txt
-rw-r--r--  2  glass           20  May 25 12:19  original.txt

```

```
% rm original.txt          ---> remove original label.

% ls -l original.txt another.txt  ---> Only another.txt exists
-rw-r--r--    1    glass          20   May 25    12:19  another.txt

% cat  another.txt          ---> list contents via other label.
  this is a file.
  hi

% _
```

# Symbolic links

- A symbolic link is an indirect pointer to a file – a pointer to the hard link to the file
  - A symlink is a special file containing a pathname
  - Like a shortcut in windows systems
- You can create a symbolic link to a directory
- A symbolic link can point to a file on a different file system
- A symbolic link can point to a non-existent file (aka. broken link)

```
#include <unistd.h>

int symlink(const char *target, const char
*linkpath);
```

# Deleting a file: unlink()

```
int unlink( const char* fileName )
```

- unlink() removes a directory entry pointing to a file.
  - If *fileName* is the last hard link to the file, the file's resources are deallocated.
  - If *fileName* is a symbolic link to the file, the link is removed.
- If successful, unlink() returns a value of 0;
- otherwise, it returns a value of -1.

# Changing a file's owner and/or group

```
int chown(const char* fileName, uid_t ownerId, gid_t groupId)
int lchown(const char* fileName, uid_t ownerId, gid_t groupId)
int fchown(int fd, uid_t ownerId, gid_t groupId )
```

- `chown()` causes **the owner and group IDs of fileName** to be changed to `ownerId` and `groupId`, respectively.
  - **A value of -1** in a particular field means that its associated value should remain unchanged.
- `lchown()`: changes the ownership of **a symbolic link**
- `fchown()`: takes **an open descriptor** as an argument instead of a filename.

Example, changed the group of the file “test.txt” from “music” to “cs” (group ID = 62).

%cat mychown.c ---> list the program.

```
main() {  
    int flag;  
    flag = chown("test.txt", -1, 62 ); /* user is ID unchanged */  
    if ( flag == -1 ) perror("mychown.c");  
}
```

% ls -lg test.txt ---> examine file before the change.

```
-rw-r--r--    1    glass    music      3  May 25 11:42    test.txt
```

% mychown ---> run program.

% ls -lg test.txt ---> examine file after the change.

```
-rw-r--r--    1    glass     cs        3  May 25 11:42    test.txt
```

# Changing file permissions

```
int chmod( const char* fileName, int mode )
```

```
int fchmod( int fd, mode_t mode );
```

- chmod() changes **the mode of *fileName*** to *mode*
- **mode** is usually supplied as an octal number,
- to change a file's mode, you must either own it or be a super-user.
- fchmod() takes **an open file descriptor** as an argument instead of a filename.
- both return **a value of -1** if unsuccessful, and **a value of 0** otherwise.

Example: changing the permission flags of the file “test.txt” to 600 octal (read and write permission for the owner only):

```
% cat mychmod.c      ---> list the program.
main() {
    int flag;
    flag = chmod("test.txt", 0600); /* Use an octal encoding */
    if ( flag==-1 ) perror("mychmod.c");
}
```

```
% ls -l test.txt      ---> examine file before the change.
-rw-r--r--      1      glass              3      May   25   11:42      test.txt
```

```
% mychmod             ---> run the program.
```

```
% ls -l test.txt      ---> examine file after the change.
-rw-----      1      glass              3      May   25   11:42      test.txt
```



# CS 288 Intensive Programming in Linux

Professor Ding, Xiaoning

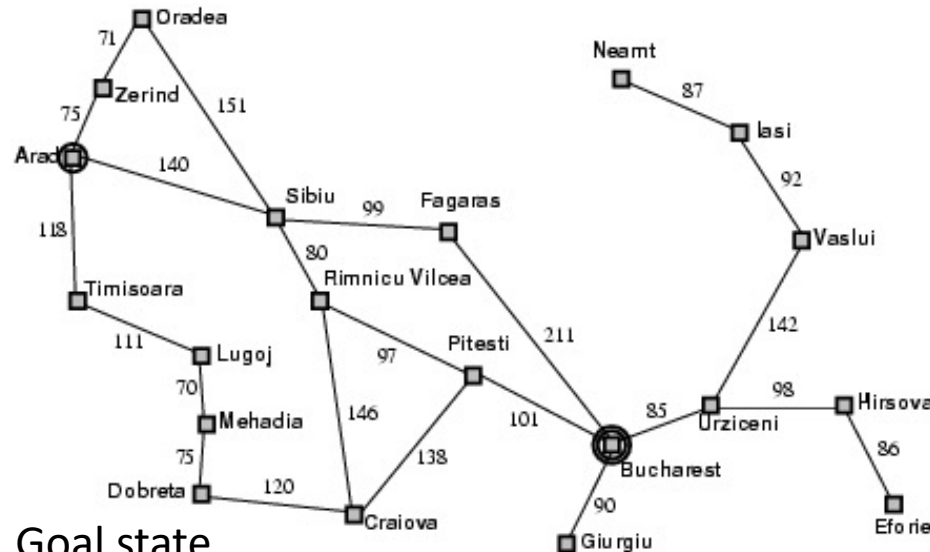

**This content may NOT be uploaded, shared, or distributed, as it is protected.**

# STATE SPACE SEARCH AND A\* SEARCH

# Solving problems by searching

- The solution is a fixed sequence of actions
- Search is the process of looking for the sequence of actions that reaches the goal

Start state



Goal state

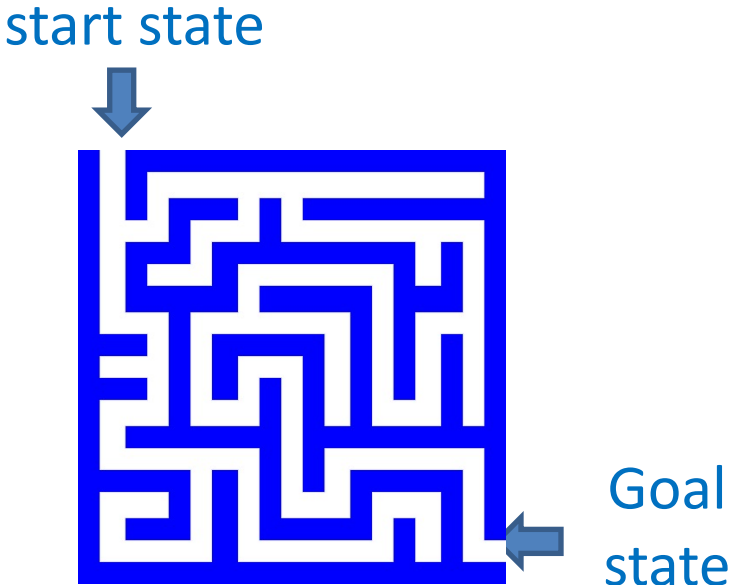
|   |   |   |
|---|---|---|
| 7 | 2 | 4 |
| 5 |   | 6 |
| 8 | 3 | 1 |

Start State

|   |   |   |
|---|---|---|
|   | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

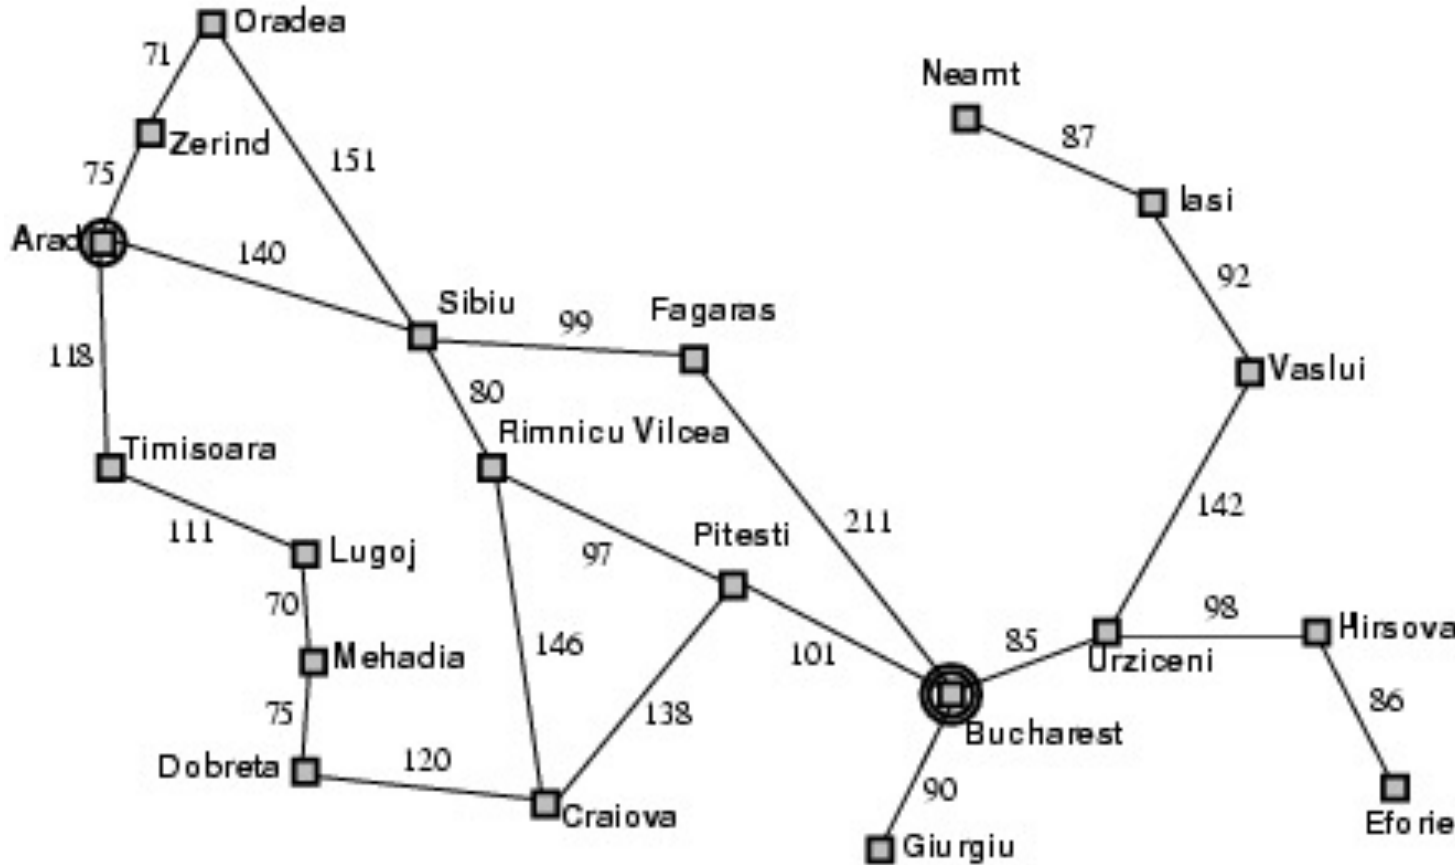
Goal State

# factors in a search problem

- **start state**
  - **Actions**
  - **Transition model**
    - What is the result of performing a given action in a given state?
  - **Goal state**
- 
- **A solution:** a sequence of actions that can change start state into goal state
  - **Path cost:** total *cost* of a sequence of actions
  - The **optimal solution** is the sequence of actions that have the lowest path cost (a problem may have multiple solutions)
  - **Search algorithms:** how do we find high quality solutions?
    - How to find a solution if there are solutions?
    - How to find a low-cost solution?
    - How to find a low-cost solution with low cost?

# Example: path finding

Find a path from *Arad* to *Bucharest*



- **Start state :** Arad
- **Actions:** Go from one city to another
- **Transition model:** If you go from city A to city B, you end up in city B
- **Goal state:** Bucharest
- **Path cost :** Sum of distances

# Example: the 8-puzzle

- **States:** Locations of tiles
  - 8-puzzle: 181,440 states
  - 15-puzzle: 1.3 trillion states
  - 24-puzzle:  $10^{25}$  states
  - Optimal solution of n-Puzzle is NP-hard
- **Transition model**
  - The movement of the empty tile.
- **Actions**
  - Move empty tile left, right, up, down
- **Path cost**
  - 1 per move  $\rightarrow$  # of moves

|   |   |   |
|---|---|---|
| 7 | 2 | 4 |
| 5 |   | 6 |
| 8 | 3 | 1 |

Start State

|   |   |   |
|---|---|---|
|   | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

# Other real-world examples

- Routing
- Touring
- VLSI layout
- Assembly sequencing
- Protein design

# Basic idea on searching for a solution

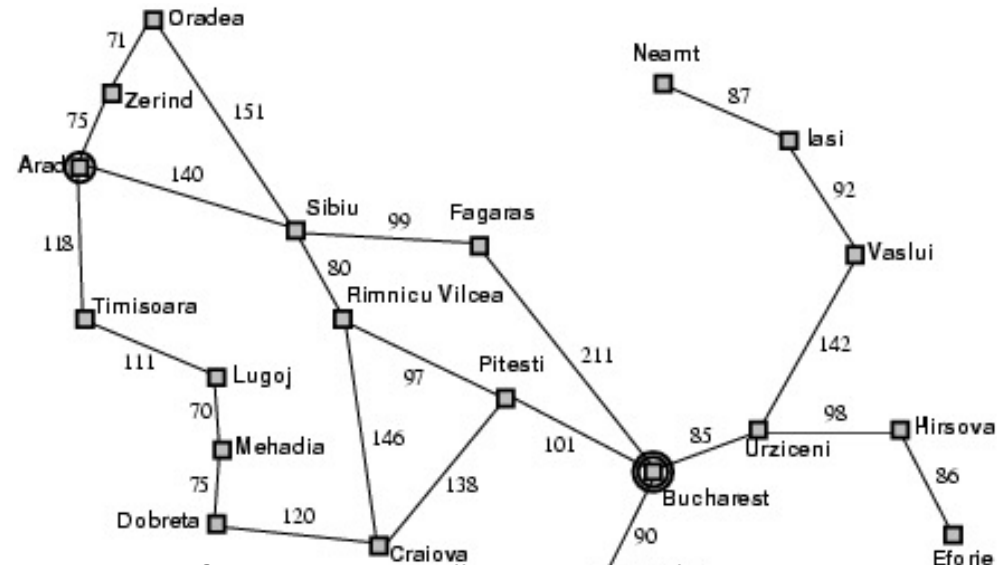

- Start from the **start state**
- Try different possible **actions**, and obtain **new states**
- Try different possible actions on new states, and obtain even more new states
- **Keep trying** until we find the goal state
- Check what actions were made to change the start state into the goal state → **solution**
- **Searching** is actually the process of trying different actions and checking new states.





# Search algorithm outline

- The set of new states is called the **fringe**.
  - key data structure --- fringe : usually implemented as a **list**, and each state is a **node** on the list.
- the fringe = {**start state**}
- While the fringe is not empty
  - **Choose** a state and remove it from fringe
  - **Check** the state: if the node contains the **goal state**, return solution.
  - **Expand** the state: for each possible action, generate a successor (new?) state, add successor states into the fringe.


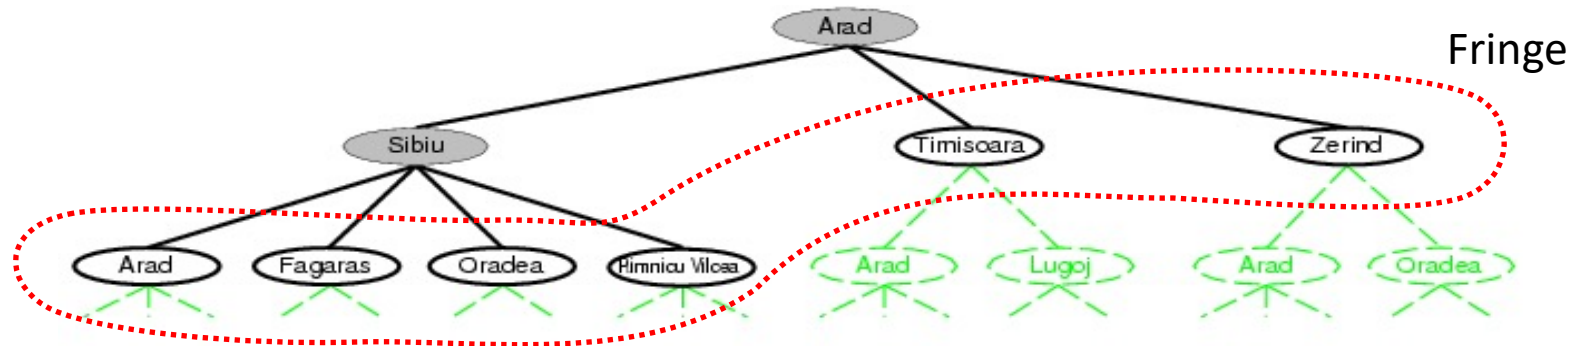
# A search example



# A search example




# A search example



# Conceptually searching a solution is to traverse a tree of states

- Root node corresponds to the start state
- Children of a node correspond to its **successor states**
- Edges correspond to actions
  - Depth of a node is how many actions were made to generate the corresponding state
  - A solution is the path starting from root and ending in a goal state
  - There may be multiple goal states. But the search can end when any one is reached.



Note: Implementation doesn't need to actually build the tree structure.

# Search strategy

**Search strategy** --- which state in the fringe should be chosen to check/expand first?

- Determine whether a solution can be found.
- Determines quality of the solution.
- Determines overhead (execution time and memory space used to find a solution).

# Uninformed and informed search strategies

**Uninformed** strategies consider all states are equally desirable.


- use only the information available in the problem definition.
- Breadth-first search
- Depth-first search
- Iterative deepening search

**Informed strategies** give the algorithm “hints” about the desirability of different states

- Use an *evaluation function* to rank states and select the most promising one for expansion
- Greedy best-first search
- A\* search

# Breadth-first search strategy


- Expand the shallowest unexpanded state
  - Depth : how many actions were made to generate the corresponding state
- Non-recursive implementation:
  - *fringe* is a FIFO queue, i.e., new successors go at end






# Breadth-first search

- Expand the shallowest unexpanded state
  - Depth : how many actions were made to generate the corresponding state
- Non-recursive implementation:
  - *fringe* is a FIFO queue, i.e., new successors go at end




# Breadth-first search

- Expand the shallowest unexpanded state
  - Depth : how many actions were made to generate the corresponding state
- Non-recursive implementation:
  - *fringe* is a FIFO queue, i.e., new successors go at end




# Breadth-first search

- Expand the shallowest unexpanded state
  - Depth : how many actions were made to generate the corresponding state
- Non-recursive implementation:
  - *fringe* is a FIFO queue, i.e., new successors go at end



# Breadth-first search

- Expand the shallowest unexpanded state
  - Depth : how many actions were made to generate the corresponding state
- Non-recursive implementation:
  - *fringe* is a FIFO queue, i.e., new successors go at end




# Pseudo-code for BFS

1. Generate an initially empty **queue** and call it OPEN.  
**OPEN is the fringe.**
2. Insert the start state into OPEN.
3. **Dequeue** a state  $n$  from OPEN. If dequeue fails because OPEN is empty, search fails (END).
4. If  $n$  is a goal state, the search succeeds (END).
5. Generate all successors to  $n$  and enqueue them into OPEN (**end of OPEN**).
6. Destroy node  $n$
7. Return to step 3.

Nodes are structures that organize states. Each node contains all the information associated to a state (e.g., state itself, cost, and pointers to link it to the fringe, etc).


# Depth-first search

- Expand deepest unexpanded state
- Implementation:
  - *fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

- Expand deepest unexpanded state
- Implementation:
  - *fringe* = LIFO queue, i.e., put successors at front



# Depth-first search


- Expand deepest unexpanded state
- Implementation:
  - *fringe* = LIFO queue, i.e., put successors at front






# Depth-first search

- Expand deepest unexpanded state
- Implementation:
  - *fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

- Expand deepest unexpanded state
- Implementation:
  - *fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

- Expand deepest unexpanded state
- Implementation:
  - *fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

- Expand deepest unexpanded state
- Implementation:
  - *fringe* = LIFO queue, i.e., put successors at front




# Depth-first search

- Expand deepest unexpanded state
- Implementation:
  - *fringe* = LIFO queue, i.e., put successors at front



# Depth-first search

- Expand deepest unexpanded state
- Implementation:
  - *fringe* = LIFO queue, i.e., put successors at front



# Pseudo-code for DFS

1. Generate an initially empty **stack** and call it OPEN.
2. Insert the start state into OPEN.
3. **Pop** a state  $n$  from OPEN. If pop fails because OPEN is empty, search fails (END).
4. If  $n$  is a goal state, the search succeeds (END).
5. Generate all successors to  $n$  and push them onto OPEN (**beginning of OPEN**).
6. Destroy node  $n$ .
7. Return to step 3.

# BFS vs DFS

- BFS
  - can always find the best solution (shortest path)
    - Always check the shallowest state (i.e., state with the lowest path cost).
  - Use a lot of memory
- DFS
  - May not find the best solution
    - Always check the deepest state. An unchecked state in fringe may be goal state. It is not selected because it is not deepest, or it has lower path cost.
  - Use less memory than BFS
  - may fail in infinite-depth paths (e.g., moving a tile back and forth)



# Basic idea on searching for a solution

- Start from the **start state**
- Try different possible **actions**, and obtain **new states**
- Try different possible actions on new states, and obtain even **more new states**
- **Keep trying** until we find the goal state
- Check what actions were made to change the start state into the goal state → **solution**
- **Searching** is actually the process of trying different actions and checking **new states**.

# DFS with CLOSED list

1. Generate an initially empty stack and call it OPEN, and a list called CLOSED.
2. Insert the start state into OPEN.
3. Pop a state  $n$  from OPEN. If pop fails because OPEN is empty, search fails (END).
4. If  $n$  is a goal state, the search succeeds (END).
5. Generate all successors to  $n$ .
6. Remove the successors that are already in OPEN or CLOSED list
7. Push the remaining successors onto OPEN (beginning of OPEN).
8. Add  $n$  into CLOSED list
9. Return to step 3.


# Informed search

- Idea: give the algorithm “hints” about the desirability of different states
  - Use an *evaluation function* to rank nodes and select the most promising one for expansion
- Greedy best-first search
- A\* search

# Heuristic function

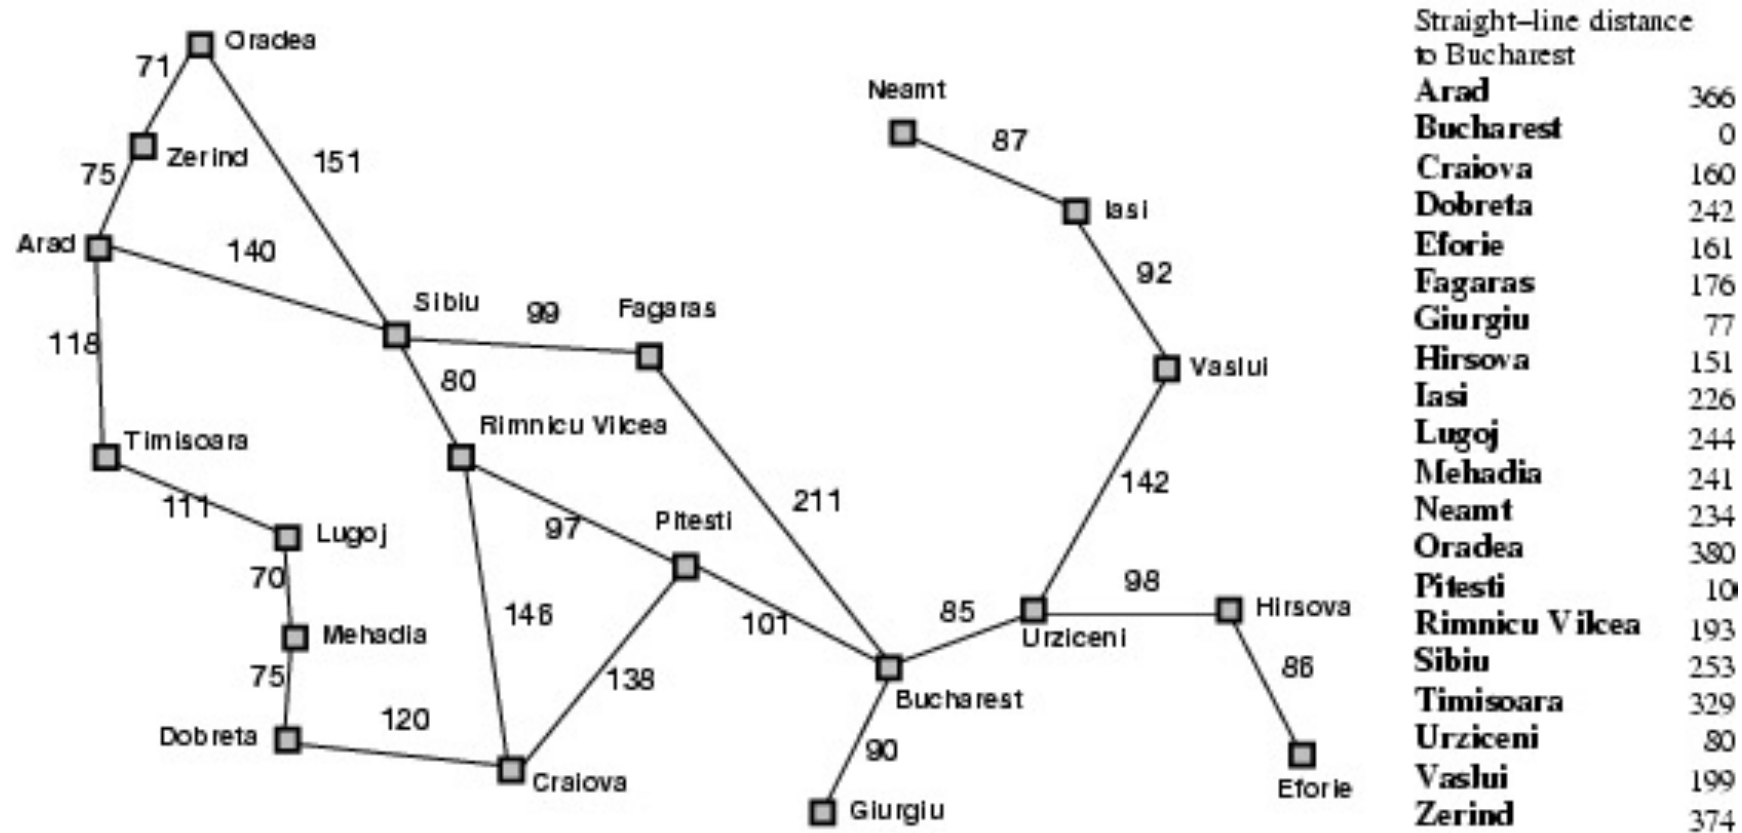
- **Heuristic function**  $h(n)$  estimates the cost of reaching goal from node  $n$
- Example:

Start state



Goal state

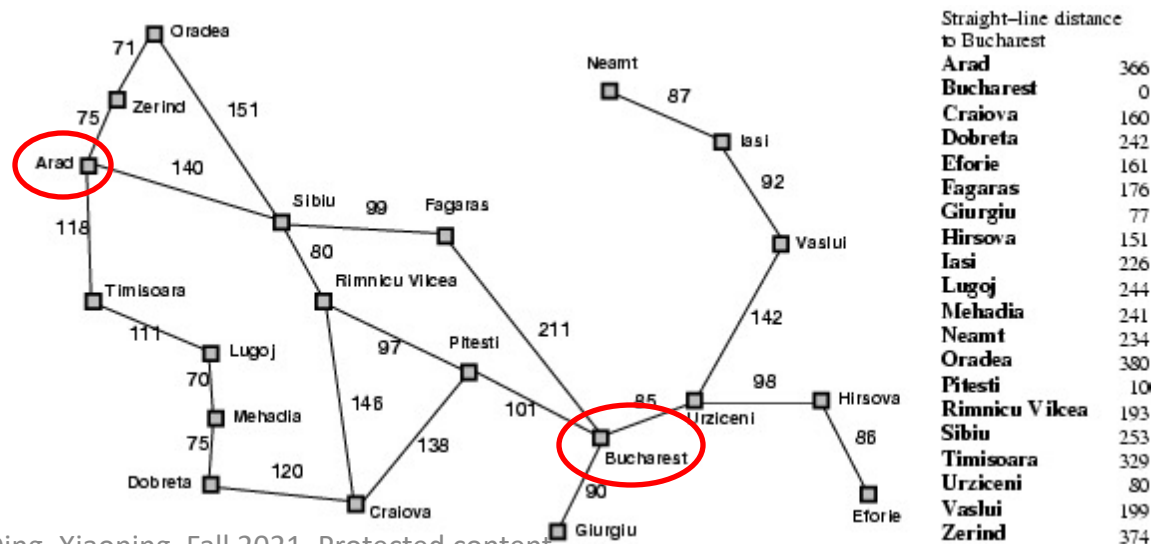
# Heuristic for the path finding problem




# Greedy best-first search

- Expand the node that has the lowest value of the heuristic function  $h(n)$

# Greedy best-first search example



# Greedy best-first search example





Straight-line distance  
to Bucharest

|                |     |
|----------------|-----|
| Arad           | 366 |
| Bucharest      | 0   |
| Craiova        | 160 |
| Dobreta        | 242 |
| Eforie         | 161 |
| Fagaras        | 176 |
| Giurgiu        | 77  |
| Hirsova        | 151 |
| Iasi           | 226 |
| Lugoj          | 244 |
| Mehadia        | 241 |
| Neamt          | 234 |
| Oradea         | 380 |
| Pitesti        | 10  |
| Rimnicu Vilcea | 193 |
| Sibiu          | 253 |
| Timisoara      | 329 |
| Urziceni       | 80  |
| Vaslui         | 199 |
| Zerind         | 374 |



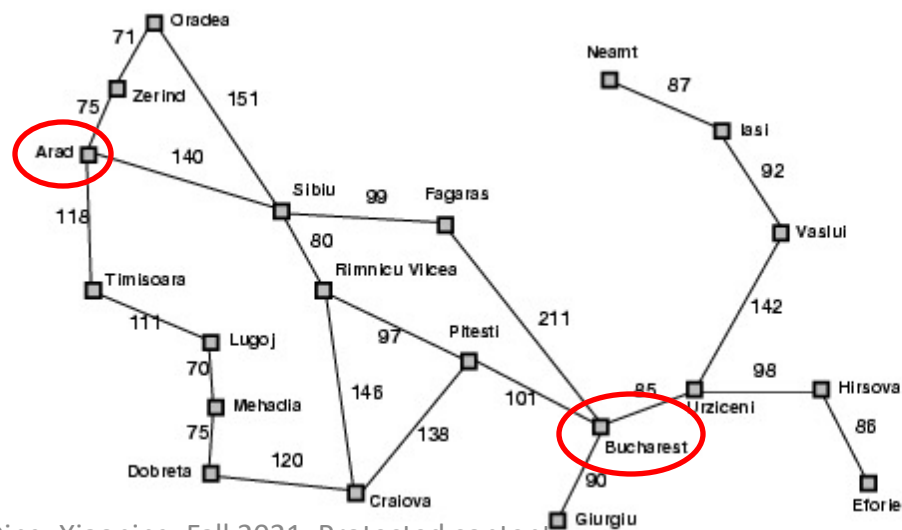
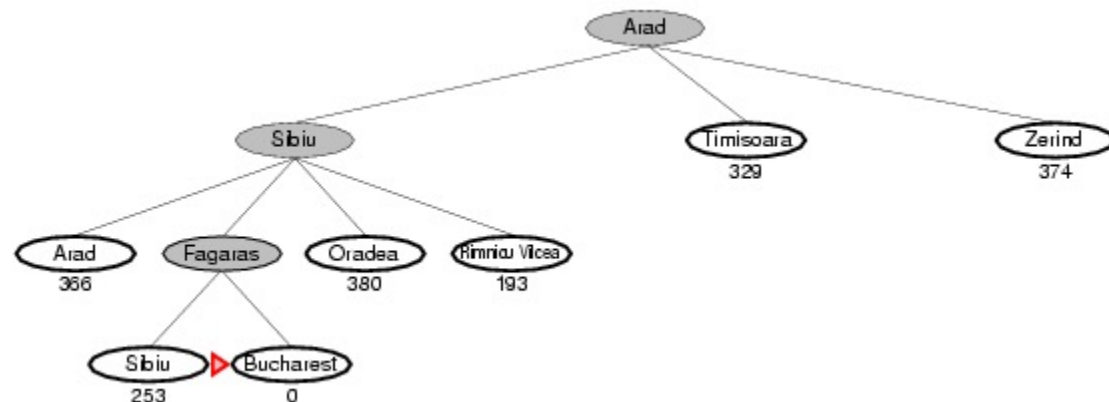
# Greedy best-first search example



Straight-line distance to Bucharest

|                |     |
|----------------|-----|
| Arad           | 366 |
| Bucharest      | 0   |
| Craiova        | 160 |
| Dobreta        | 242 |
| Eforie         | 161 |
| Fagaras        | 176 |
| Giurgiu        | 77  |
| Hirsova        | 151 |
| Iasi           | 226 |
| Lugoj          | 244 |
| Mehadia        | 241 |
| Neamt          | 234 |
| Oradea         | 380 |
| Pitesti        | 10  |
| Rimnicu Vilcea | 193 |
| Sibiu          | 253 |
| Timisoara      | 329 |
| Urziceni       | 80  |
| Vaslui         | 199 |
| Zerind         | 374 |


# Greedy best-first search example



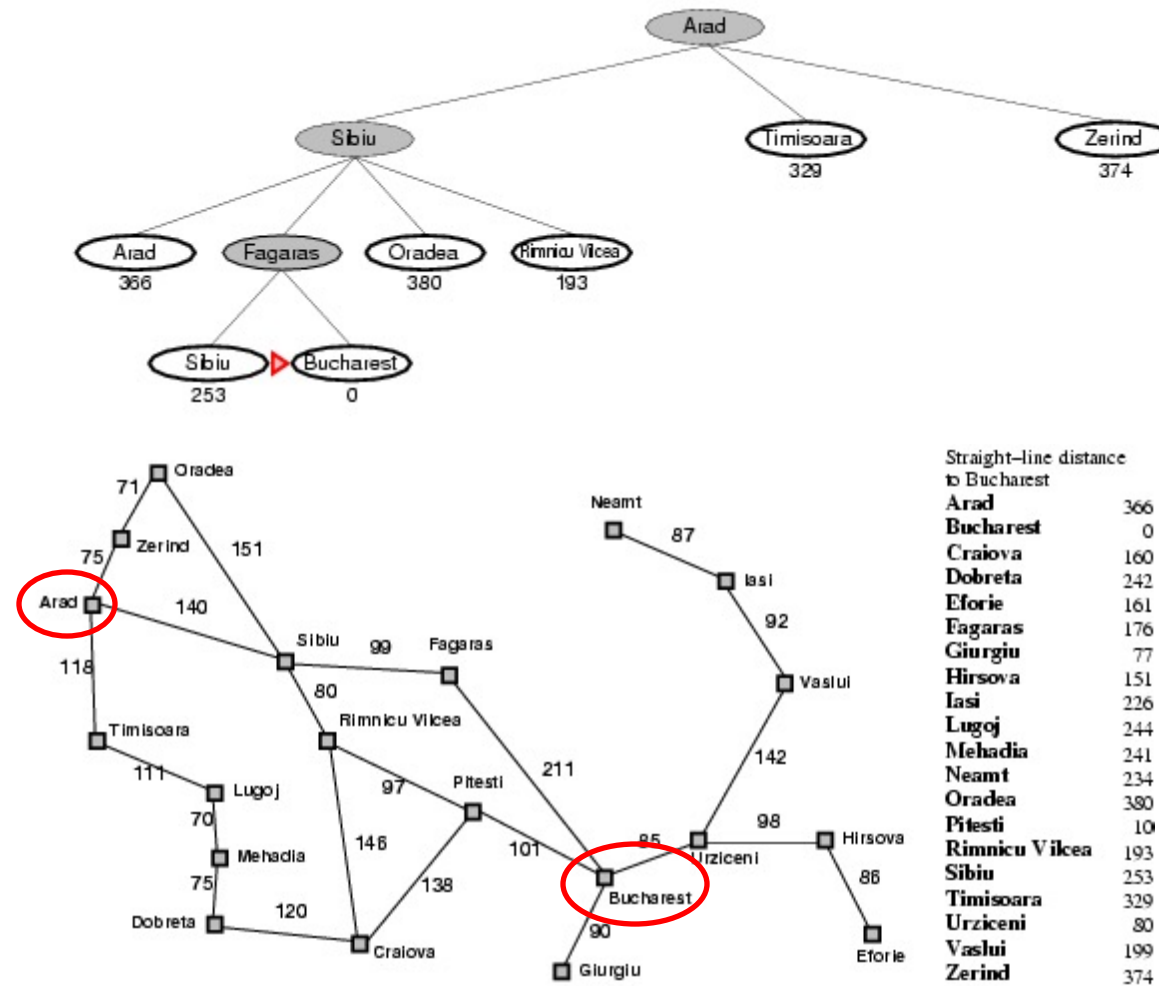
Straight-line distance to Bucharest

|                |     |
|----------------|-----|
| Arad           | 366 |
| Bucharest      | 0   |
| Craiova        | 160 |
| Dobreta        | 242 |
| Eforie         | 161 |
| Fagaras        | 176 |
| Giurgiu        | 77  |
| Hirsova        | 151 |
| Iasi           | 226 |
| Lugoj          | 244 |
| Mehadia        | 241 |
| Neamt          | 234 |
| Oradea         | 380 |
| Pitesti        | 10  |
| Rimnicu Vilcea | 193 |
| Sibiu          | 253 |
| Timisoara      | 329 |
| Urziceni       | 80  |
| Vaslui         | 199 |
| Zerind         | 374 |

# Greedy best-first search may get stuck in loops



# Greedy best-first search is not optimal



# A\* search

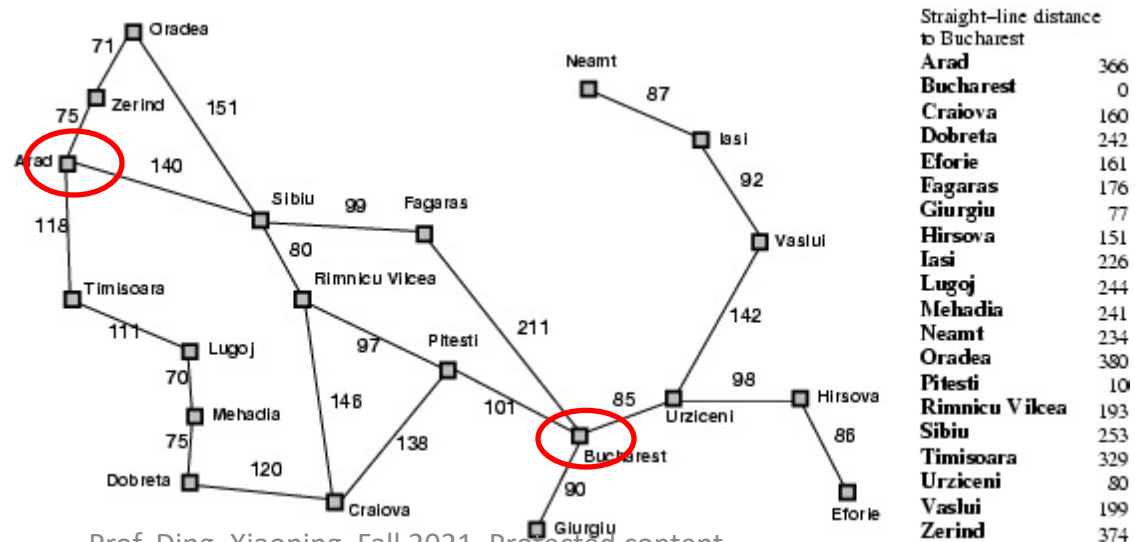

- Idea: avoid expanding paths that are already expensive
- The evaluation function  $f(n)$  is the estimated total cost of the path through node  $n$  to the goal:

$$f(n) = g(n) + h(n)$$


$g(n)$ : cost so far to reach  $n$  (path cost)

$h(n)$ : estimated cost from  $n$  to goal (heuristic)

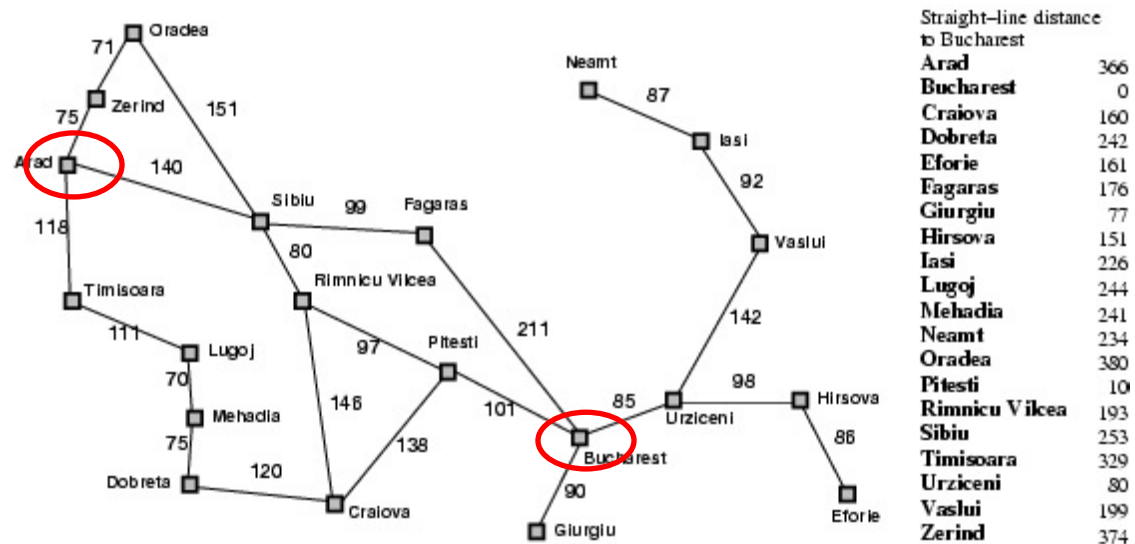

# A\* search example



# A\* search example


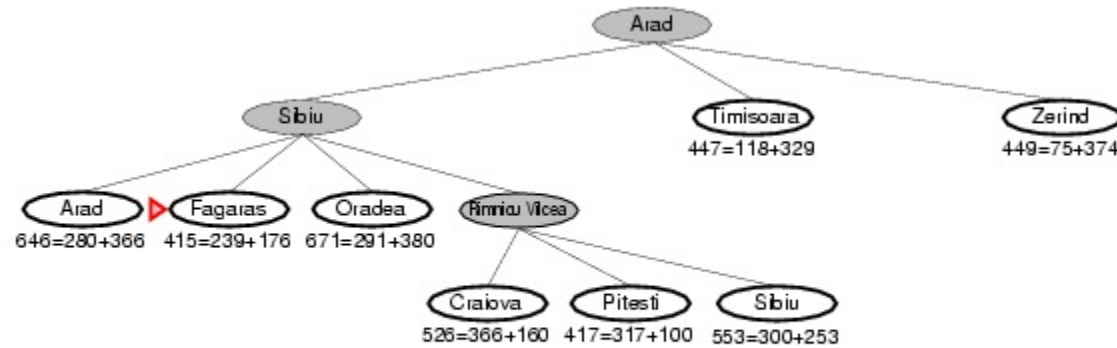


# A\* search example








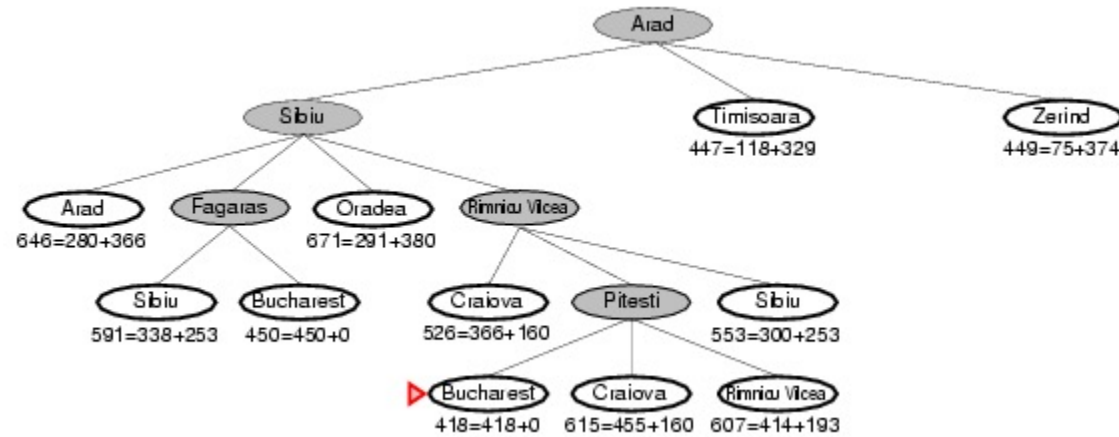
# A\* search example



# A\* search example



# A\* search example



Straight-line distance to Bucharest

|                |     |
|----------------|-----|
| Arad           | 366 |
| Bucharest      | 0   |
| Craiova        | 160 |
| Dobreta        | 242 |
| Eforie         | 161 |
| Fagaras        | 176 |
| Giurgiu        | 77  |
| Hirsova        | 151 |
| Iasi           | 226 |
| Lugoj          | 244 |
| Mehadia        | 241 |
| Neamt          | 234 |
| Oradea         | 380 |
| Pitesti        | 10  |
| Rimnicu Vilcea | 193 |
| Sibiu          | 253 |
| Timisoara      | 329 |
| Urziceni       | 80  |
| Vaslui         | 199 |
| Zerind         | 374 |

# Pseudo-code for A\*

1. Generate an initially empty **priority** queue and call it OPEN, **prioritized by  $f(n)$  value**.
2. Generate an initially empty collection of states and call it CLOSED.
3. Insert the start state (or states) into OPEN.
4. Remove the state with smallest  $f(n)$  value (state  $n$ ) from OPEN. If removal fails because OPEN is empty, search fails (END).
5. If  $n$  is a goal state, finish the search.
6. Generate all successors to  $n$ ;
7. **Remove and free the successors that are already in OPEN or CLOSED list.**
8. Push remaining successors onto OPEN (**prioritized by  $f(n)$  value**).
9. Add  $n$  to CLOSED.
10. Return to step 4.

# Admissible heuristics

- A heuristic  $h(n)$  is **admissible** if for every node  $n$ ,  $h(n) \leq h^*(n)$ , where  $h^*(n)$  is the true cost to reach the goal state from  $n$
- An admissible heuristic never overestimates the cost to reach the goal, i.e., it is optimistic
- Example: straight line distance never overestimates the actual road distance
- Theorem: If  $h(n)$  is admissible,  $A^*$  is optimal

# Designing heuristic functions

- Heuristics for the 8-puzzle

$h_1(n)$  = number of misplaced tiles

$h_2(n)$  = total Manhattan distance (number of squares from desired location of each tile)

|   |   |   |
|---|---|---|
| 7 | 2 | 4 |
| 5 |   | 6 |
| 8 | 3 | 1 |

State being evaluated

|   |   |   |
|---|---|---|
|   | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

$$h_1(\text{start}) = 8$$

$$h_2(\text{start}) = 3+1+2+2+2+3+3+2 = 18$$

- Are  $h_1$  and  $h_2$  admissible?

# 15-puzzle problem

